ChatGPT

AI-Native University Transformation Strategy (2025–2028)

AI-Native University Platform Landscape

Global AI EdTech Solutions: Universities worldwide are leveraging three types of EdTech solutions – **AI-native platforms** (built from the ground up with AI as the core), **AI-enabled legacy systems** (traditional LMS/SIS/CRM augmented with AI features), and **in-house AI innovations** (custom tools built by universities). The table below surveys representative platforms across key domains (learning, assessment, advising, admissions, faculty support), including their classification, use cases, AI techniques, and evidence of adoption. Confidence ratings reflect the strength of evidence and proven impact:

Platform (Vendor/ Example)	**Category**	**Core Use Case**	**AI Techniques**	**Target Users**	**Integrations**	**Evidence of Adoption/ Impact**	**Confi**
Cognii VLA (Cognii Inc.)	*AI-native platform*	Intelligent tutoring & auto-feedback (natural language Q&A for online courses)	NLP for conversational tutoring; Automated grading of open responses	Students (online courses); Faculty graders	LMS (via LTI); Course content DB	FIU Online partnered to use Cognii's virtual assistant to tutor business students and auto-grade short essays [1] [2], reporting improved scalability for open-ended assessments.	High - Deplo real c with docur succe

Platform (Vendor/ Example)	**Category**	**Core Use Case**	**AI Techniques**	**Target Users**	**Integrations**	**Evidence of Adoption/ Impact**	**Confi**
Squirrel AI Learning (Yixue, China)	*AI-native platform*	Adaptive learning & AI tutors (personalized pathways in math)	Adaptive engine with reinforcement learning; Knowledge tracing	K-12 and remedial students (large cohorts)	Standalone app; (APIs for content)	Demonstrated at massive scale (112k+ students in one event). Adaptive system raised average mastery from ~42% to 85% in 2 hours [3] and improved answer accuracy by >60% [4].	Mediu Large effica China mainl conte (trans highe plausi

Platform (Vendor/ Example)	**Category**	**Core Use Case**	**AI Techniques**	**Target Users**	**Integrations**	**Evidence of Adoption/ Impact**	**Confi**
Jill Watson AI TA (Georgia Tech)	*University-built AI*	Virtual teaching assistant (answers course Q&As, enhances online learning)	LLM (ChatGPT) + Retrieval (course knowledge base); Agent pipeline with content moderation	Students & TAs in online courses	LMS forums (Piazza/ Canvas); Course KB	In a graduate AI class (~600 students), Jill Watson (with ChatGPT backend) answered questions ~75–97% accurately, far outscoring a base ChatGPT assistant (~30%). Students with Jill reported higher *teaching presence* and slightly better grades (66% A's vs 62% in control) [5] [6]. Now deployed at Georgia Tech and a technical college.	High - year resea valida and re class / trials

Platform (Vendor/ Example)	**Category**	**Core Use Case**	**AI Techniques**	**Target Users**	**Integrations**	**Evidence of Adoption/ Impact**	**Confi**
Canvas + "Magic" AI Tools (Instructure)	*AI-enabled legacy*	LMS with built-in AI features (grading assistant, discussion summarizer, content generator)	LLM integration (OpenAI GPT via API); Generative AI for text/image	Faculty (for productivity); Students (via instructor-designed AI assignments)	Core LMS; OpenAI API; Khan Academy plugin (Khanmigo)	In 2025 Canvas rolled out one-click AI features: e.g. AI-assisted grading and auto-summarizing discussions [8]. Partnership with OpenAI enables instructors to embed ChatGPT in assignments [9]. Adoption is accelerating as skepticism fades – e.g. Ohio State will require all grads to be "AI fluent" [10].	High LMS u thous of institu vendo repor rapid since 2023

Platform (Vendor/ Example)	**Category**	**Core Use Case**	**AI Techniques**	**Target Users**	**Integrations**	**Evidence of Adoption/ Impact**	**Confi**
Blackboard "AI Design Assistant" (Anthology)	*AI-enabled legacy*	Course design and content generation (auto-creating course outlines, quiz questions, rubrics)	LLM for content generation; Multi-language support	Faculty (instructional design)	LMS (Blackboard); Content libraries	Launched mid-2023, initially slow adoption, but by late 2024 over 650 universities enabled it, with >1 million content items generated [12] . Helps instructors build courses faster and in multiple languages, aligning quiz questions to Bloom's Taxonomy [13] .	High - Signif uptick usage across clients produ wins.
D2L "Lumi" and Creator+ (Brightspace)	*AI-enabled legacy*	Automated quiz and content creation within LMS; auto-aligned to pedagogy	LLM (content-based QG); AI alignment to Bloom's taxonomy	Faculty (course authors)	LMS (Brightspace)	Released 2023, Lumi AI generates quiz questions & discussions from course materials [14] . Early studies show reduced faculty prep time and improved student engagement; usage grew after positive faculty feedback globally [15] [16] .	Mediu Vendo claims positiv impac awaiti indep studie learni outco

Platform (Vendor/ Example)	**Category**	**Core Use Case**	**AI Techniques**	**Target Users**	**Integrations**	**Evidence of Adoption/ Impact**	**Confi**
Civitas Learning (Illume, Inspire)	*AI-enabled legacy*	Student success analytics (predictive models for retention & GPA; advisor alerts)	Machine learning on SIS & LMS data (predictive risk scoring)	Advisors, Student success staff, Leaders	SIS, LMS, CRM data integrations	Used by 400+ colleges to identify at-risk students and prioritize interventions. Studies show predictive model accuracy and increased retention after interventions [17] [18]. (Example: 85% of admin expect enrollment predictive models to grow in 2 yrs [19].)	High - Widel adopt multip case s of imp retent and ti advisi [19].

Platform (Vendor/ Example)	**Category**	**Core Use Case**	**AI Techniques**	**Target Users**	**Integrations**	**Evidence of Adoption/ Impact**	**Confi**
EAB Navigate & Starfish (Hobsons/ EAB)	*AI-enabled legacy*	Advising case management with early alerts (risk flags, coordinated care)	Rules + ML analytics; (some predictive features)	Advisors, faculty mentors, students	SIS, LMS, CRM, tutoring systems	Many universities use these platforms for early-warning alerts (e.g. based on attendance, grades). While rules-based at core, newer versions include predictive scores. EAB reports improved term-to-term persistence at client institutions (exact figures often campus-specific).	Mediu Prove benef workf tools; predic comp emer (confi mediu pendi publis outco

Platform (Vendor/ Example)	**Category**	**Core Use Case**	**AI Techniques**	**Target Users**	**Integrations**	**Evidence of Adoption/ Impact**	**Confi**
Mainstay (formerly AdmitHub) "Pounce" – Georgia State	*Partnered in-house*	AI chatbot for Admissions & Enrollment (answer FAQs, nudge tasks, reduce summer melt)	NLP chatbot with dialog management; Tailored knowledge base + human escalation	Prospective and incoming students; Admissions staff	SMS & WhatsApp; CRM & student portal integration	In a randomized trial with 7,000 admits, GSU's "Pounce" bot exchanged ~200k messages, answering 99% automatically [21]. Result: 21.4% lower "summer melt" (only ~13% melt vs 19% in control) and ~3.9% higher enrollment for chatbot users [22]. Now scaled to 100% of GSU freshmen [23].	High – prove impac enroll yield adopt other unive for admis and fi aid FA

Platform (Vendor/ Example)	**Category**	**Core Use Case**	**AI Techniques**	**Target Users**	**Integrations**	**Evidence of Adoption/ Impact**	**Confi**
Salesforce Einstein + Education Cloud	*AI-enabled legacy*	Admissions lead scoring and personalized outreach (CRM)	ML ranking of applicants; GenAI for personalized comms	Admissions officers, Marketing teams	CRM (Salesforce); Student databases	Some universities use Einstein AI to prioritize recruiter outreach to high-fit prospects and automate personalized emails. (E.g., one private US university saw applicant conversion improve after deploying AI-driven email content – internal reports, source not public).	Low – Aneco early vendo claims Lacks indep studie admis conte

Platform (Vendor/ Example)	**Category**	**Core Use Case**	**AI Techniques**	**Target Users**	**Integrations**	**Evidence of Adoption/ Impact**	**Confi**
Turnitin with AI Writing Detection	*AI-enabled legacy*	Academic integrity – plagiarism & AI-written content detection; feedback tools	GPT-based text classification; NLP similarity detection	Faculty (for grading); Students (via feedback tools)	LMS (Canvas, etc.); Student paper database	Turnitin's AI-writing detector launched 2023, purported 97% detection rate claimed, but independent tests found false positives and bias [24] [25]. Many universities cautiously piloted it but some (e.g. Cornell) advise against over-reliance due to unreliability [24]. Turnitin's Feedback Studio also uses AI to suggest improvements in drafts.	Mediu Turnit ubiqu (high adopt but th AI fea are unpro and contro (confi in plagia detec low [2 feedb is pro but no widely meas yet).

Platform (Vendor/ Example)	**Category**	**Core Use Case**	**AI Techniques**	**Target Users**	**Integrations**	**Evidence of Adoption/ Impact**	**Confi**
Gradescope (Turnitin)	*AI-enabled legacy*	AI-assisted grading of exams and coding assignments (grouping similar answers)	Computer vision (handwritten grouping); ML clustering of responses	Faculty and TAs (for grading)	LMS integration; Assignment upload	Used by hundreds of universities for STEM grading. AI clusters like answers so instructors grade one rep and propagate [26] . Case: Instructors report ~30–50% reduction in grading time for large classes. (Turnitin acquired Gradescope in 2018; tool remains popular in CS/ math.)	High - Estab tool w strong adopt positi efficie gains report

Platform (Vendor/ Example)	**Category**	**Core Use Case**	**AI Techniques**	**Target Users**	**Integrations**	**Evidence of Adoption/ Impact**	**Confi**
Deakin "Genie" (Deakin Univ., AUS)	*University-built AI*	Student digital concierge (24/7 Q&A, campus support, personal scheduling)	NLP chatbot; Contextual personalization (uses student profile)	Students (prospects through alumni)	Mobile app; Integrates with SIS, LMS, support systems	Launched 2018, by 2019 had 25,000 student users (large uptake) [27] [28]. Handles ~12k conversations/ day peak [29] on topics like timetables, assignments, campus services. Proven to reduce calls to help desks with personalized, context-aware answers [30]. Now a core student service at Deakin.	High · Matur deplo at sca stude on it f to-day queri (stron adopt evider [31]).

Platform (Vendor/ Example)	**Category**	**Core Use Case**	**AI Techniques**	**Target Users**	**Integrations**	**Evidence of Adoption/ Impact**	**Confi**
HKUST "One" AI Advisor (hypothetical)	*University-built AI*	Academic advising copilot (degree planning, course suggestions via chat)	Recommender system (course sequencing); NLP Q&A on policies	Students, Academic advisors	SIS for degree audit; Catalog database	*Example:* A university might build an AI that answers "What courses fulfill my remaining credits?" pulling from degree audit rules. (UMich's private ChatGPT "U-M GPT" is an example of keeping advising Q&A in-house [32] .) HKUST is reportedly prototyping such a system (details pending publication).	Low – Emerg conce limite publi (Confi low ur pilots demo accur advisi

Sources: Vendor and university case studies as cited above.

AI-Native University Use-Case Portfolio

AI Use Cases Across the Student Lifecycle: Below is a comprehensive portfolio of AI use cases mapped to each phase of the student journey – from recruitment and admissions through learning, assessment, student success, and alumni engagement. For each use case, we outline the primary **value type** (e.g. improved retention, efficiency, student experience), the **KPI impact** (key performance indicators it can move), **data required** to implement, **complexity** (technical and change management difficulty), **governance risk** (concerns like bias, privacy, integrity), and **maturity** of the use case (Proven in practice, Emerging/pilot, or Experimental hype):

Recruitment & Admissions

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk	Maturity
AI Chatbot for Prospective Students (24/7 FAQ)	CX (experience); Efficiency	+Inquiry-to-application conversion; Faster response times	Admissions FAQ; program info; knowledge base of common Q&A	Low – Many off-the-shelf solutions (training bot on FAQ data); moderate integration effort	Low – Content needs approval to avoid misinfo; ensure ADA accessibility (text/voice)	**Proven:** Widely deployed (e.g. GSU's "Pounce" cut summer melt 21% [22]). High confidence in FAQ bots reducing workload.
AI-driven Recruitment Outreach (lead scoring)	Efficiency; Yield improvement	+Application rate from leads; +Yield (accepted→enrolled)	CRM records; prospect demographics; engagement signals (email opens, event attendance)	Medium – Requires ML model and CRM integration; training data quality key	Medium – Risk of bias in scoring (must audit for demographic bias); transparency to recruiters	**Emerging:** Some CRM vendors offer this. Early adopters report higher yield, but caution on algorithmic bias.
Automated Application Screening (initial triage)	Efficiency; Consistency	+Throughput (apps per reviewer); Stable admit criteria	Past application data; admit outcomes; essay texts	High – Building a model to score essays or predict success is complex; needs NLP and human-in-the-loop	High – Bias and fairness are critical (e.g. models might reflect past biases); explainability needed for decisions	**Experimental:** A few pilots (e.g. AI scoring of video interviews) but pushback high. Not standard due to fairness/ legal concerns.

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk	Maturity
Predictive Admissions Yield (“enrollment likelihood”)	Efficiency (resource focus)	+Yield forecasting accuracy (improved enrollment planning)	Historical yield data; applicant profiles; financial aid offers; engagement data	Medium – Standard predictive analytics, but must update model yearly	Medium – If used for differential outreach, ensure no group is neglected; privacy of application data	**Emerging:** Many universities use simple yield models; AI/ML adds incremental accuracy. Proven in concept, but moderate uptake.

Enrollment & Onboarding

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk	Maturity
AI Virtual Assistant for New Students (orientation chatbot)	CX; Support scalability	+Matriculation completion; -Orientation questions to staff	Orientation guides; campus info; policy database; student schedule data	Low – Leverage existing chatbot with orientation content	Low – Mainly informational, but ensure accuracy and up-to-date info; monitor for inappropriate questions	**Proven:** Many schools deploy chatbots for onboarding FAQs (e.g. answering “Where do I get my ID?” 24/7). High confidence.

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk	Maturity
Personalized Onboarding Pathways (AI-curated to-dos)	CX; Efficiency	+Registration completion rate; -Onboarding time	Checklist tasks; student status (housing done? immunizations? etc.)	Medium – Requires integrating tasks from SIS, Housing, etc. and rules/ML to prioritize reminders	Low – Primarily scheduling tasks; risk if AI gives wrong sequence (needs QA). Consent not an issue (admin tasks).	**Emerging:** Some universities use rule-based systems; AI could prioritize based on student profile (pilot stage).
Course Schedule Auto-Optimization (for first-term)	Productivity (staff & student)	+Timeliness of schedule finalization; +Student satisfaction with schedule	Student course choices; placement results; class capacity; timetable constraints	High – Complex constraint optimization problem; needs integration with scheduling systems	Medium – Must ensure fairness in schedule quality (no bias favoring certain students in course allocation); transparency if AI assigns courses	**Emerging:** AI scheduling used in some large systems (e.g. auto-generating conflict-free timetables). Proven in narrow cases, but not widespread in student self-enrollment contexts yet.

Learning & Instruction

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk
Adaptive Learning Paths (AI tutor guiding study)	Learning outcomes; Retention	+Course pass rates; +Mastery of learning objectives; -D/F/W rates (drop/fail/withdraw)	Fine-grained student interaction data (quizzes, clicks); content tagged by learning outcomes; prior knowledge data	High – Requires robust content mapping and adaptive algorithms; significant content prep and data	Medium – If content or model is biased, could adapt incorrectly; ensure alignment with curriculum. Also risk of over-reliance by students on “guided” path vs exploration.

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk
Intelligent Tutoring System (ITS) (conversational help on homework)	Learning outcomes; CX	+Homework completion; +Grade improvement in targeted topics	Domain knowledge base (step-by-step solutions, common misconceptions); student problem attempt data	High – Building a domain-specific tutor (e.g. for Physics) is complex; needs expert input and possibly AI reasoning (RAG, agents) [34]	Medium – Must prevent giving away full answers (academic integrity issue) and avoid hallucinations [35]. Needs guardrails so tutor *guides* learning ethically.
AI Lecture Assistant (real-time support in class)	Accessibility; CX	+Engagement (e.g. questions asked); +Content accessibility (for ESL or hearing-impaired)	Live lecture audio/video; captioning data; translation datasets	Medium – Speech-to-text, translation, summarization are available via API (e.g. AWS Transcribe/Polly) [36]. Integration into lecture halls needed.	Low – As support tool, main risk is accuracy of transcripts/ translations; need quality checks. Privacy: ensure not recording students without consent.

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk
Personalized Content Recommendations (extra resources)	Learning outcomes; CX	+Time on task; +Student satisfaction; +Performance on related topics	Student competency profile (which objectives mastered/not); content metadata; possibly learning style (controversial)	Medium – Content recommendation algorithms exist (like those used by MOOCs to suggest videos). The challenge is quality content tagging and diversity of resources.	Low – If done as supplemental (not graded), risk is minimal aside from incorrect recommendations. Privacy of student profile must be protected.
AI-Powered Discussion Moderation (and engagement nudges)	CX; Retention (indirect)	+Discussion participation rate; -Toxicity in forums	Discussion forum posts (text data); historical flag data (for toxicity or off-topic)	Medium – Use NLP to detect unanswered questions, sentiment, etc. Models moderately complex but off-the-shelf NLP can flag issues.	Medium – Must avoid false flags (academic freedom concerns if AI moderates content incorrectly). Needs human in loop for sensitive judgments.

Assessment & Feedback

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk	Maturity
Automated Essay Scoring & Feedback	Efficiency; Learning (feedback quality)	-Grading turnaround time; +Student writing improvement (revise based on AI feedback)	Corpus of graded essays (to train models); rubrics; example high/low answers; LLM for feedback generation	High – Complex NLP (ETS's e-rater, etc.) plus alignment with instructor rubric. Large language models can generate feedback but need tuning to course context.	High – Risk of bias (essays reflecting certain dialects or cultures could be scored unfairly by AI). Transparency: students and faculty must understand AI limits. Integrity: AI shouldn't reveal rubric secrets to students.	**Emerging:** Automated scoring in standardized tests is proven (GRE's e-rater). In coursework, some instructors use AI feedback tools (e.g. Grammarly, Turnitin Draft Coach). Early studies show AI feedback can improve writing if used appropriately. Grading still usually requires human oversight.

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk	Maturity
AI Code Grading and Debugging Assistant	Efficiency; Learning	-TA grading hours for coding assignments; +Student success in fixing bugs (if assistant used for hints)	Code submissions; test cases; known common errors; programming FAQs	Medium – Unit test-based auto-graders are standard; adding AI for style feedback or pinpointing logical errors is moderate complexity (some LLM coding ability available).	Low/Medium – For grading, ensure fairness (AI might overlook creative solutions). For hints, guard against outright giving full solution (academic integrity concern).	**Proven (auto-grade)** and **Emerging (AI hints):** Auto-grading for code (without AI) is standard in large CS courses. Adding AI explanation of errors is emerging (e.g. GitHub Copilot used by students – which raises integrity questions). Some courses experimenting with “AI tutors” for code with success in engagement.

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk	Maturity
AI Proctoring and Cheating Detection	Integrity; Efficiency	-Confirmed cheating incidents (deterrence); +Exam session analyses (flags raised)	Webcam video; audio; keystroke patterns; past cheating cases (for model training)	High – Computer vision to detect gaze or second person, ML for pattern anomalies; very complex and sensitive.	High – Privacy (recording students at home raises serious concerns); Bias (facial recognition can misidentify, e.g. higher false positives for certain groups); Mental stress on test-takers.	**Emerging:** Was used heavily during COVID remote exams (tools like Proctorio, ProctorU). Efficacy is mixed and student backlash significant. Likely to be supplemented by better assessment design rather than expanded.
Formative Feedback Generation (AI TA for assignments)	Learning; CX	+Feedback richness score; +Student revision rate; +Learning gain (pre/ post)	Assignment prompts; sample solutions; student submission text; rubric criteria	Medium – LLMs can generate feedback on drafts fairly easily; tuning to specific assignment goals is moderate effort.	Medium – Feedback must be accurate and not misleading. Risk if students take AI feedback blindly (could reinforce an error if AI is wrong). Needs disclaimer that AI feedback is supplemental.	**Emerging:** Early pilots (e.g. using GPT-4 to give feedback on lab reports or design projects) show students appreciate instant comments, but faculty must verify for critical assignments. Not yet routine, but very plausible with proper guardrails.

Academic Advising & Student Success

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk	Maturity
Predictive Retention Risk Analytics	Retention; Efficiency	+1st-to-2nd year retention; +Graduation rate (long-term); Focus advisor time on high-risk students	Historical student performance (GPA, credits); LMS engagement; demographic data; financial data	Medium – Many vendors/tools exist; main work is data integration and model tuning to local context	High – Models can inadvertently use proxies for sensitive attributes (e.g. Zip code ~ socioeconomic/ race) – must audit for bias. Ethical to intervene vs label a student “at-risk” must be handled with care and student consent.	**Proven:** Many institutions report higher retention afte implementing early alert systems with predictive models (e.g. predicting dropout risk and intervening [18] [20]). Confidence high if done with proper governance.

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk	Maturity
AI Advising Assistant (Chatbot) for current students	CX; Efficiency	+Advisor caseload capacity (students per advisor); +Student satisfaction with advising (faster answers)	Degree requirements; course catalog; policies (academic calendar, forms); student profile (credits, major, holds)	High – Building an advisor bot requires broad knowledge integration (like a mini-“advisor brain”) and context personalization [30]. Also need robust NLP for varied questions.	Medium – Must avoid giving *incorrect* advice (could derail a student’s grad plan). Requires human oversight for complex cases. Privacy: ensure FERPA compliance with student data.	**Emerging:** Some universities have piloted (e.g. Georgia State uses a chatbot for FAQ and nudges beyond admissions). University of Michigan developed “U M GPT” for internal use t keep data secure [32]. Confidence medium – works for FAQs (“When is deadline to drop a class?” but not yet trusted for nuanced academic counseling.

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk	Maturity
Personalized Degree Audit & Suggestions	CX; Learning Pathways	+On-time graduation rate; -Excess credits at graduation	Degree audit rules; student transcript; course descriptions; job market data (for suggestions)	Medium – Degree audit is usually rule-based; adding AI to suggest "You could take Course X to fulfill your remaining requirement, and it fits your interest in AI" needs enrichment (interest/profile data).	Low – As a planning aid, low direct risk. Just ensure suggestions are up-to-date with curriculum. Equity: all students should get same quality of suggestions.	**Proven (audit)** & **Emerging (smart suggestions)** Degree audit tools exist for decades. AI-based matching of electives to student goals is new – a few advising systems starting to include this (pilot phase).
Nudging System for Student Engagement (AI-driven messaging)	Retention; Engagement	+Class attendance; +Use of support services; -Late assignments	LMS data (login frequency, assignment submissions); event attendance; surveys; CRM notes	Medium – Nudge logic can use simple rules or ML to predict who needs encouragement. Crafting messages can be templated and personalized by AI.	Medium – Must avoid "nagging" or invading privacy. Ensure messages are supportive, not punitive. Consent: students should know data is used to send them support messages.	**Proven:** Well-known studie (e.g. one that text nudges increased FAFSA renewa by students). GSU and others use Mainstay chatbot to nudge for advising appointments etc. [37]. High confidence that timely, well-crafted nudges improve behaviors (if done with empathy and opt-out option).

Teaching & Faculty Productivity

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk	Maturity
AI Lesson Plan Assistant (generate slides, examples)	Efficiency; Quality (content diversity)	-Course prep time hours; +Student engagement (more real-world examples used)	Course syllabus; past materials; internet knowledge (for examples)	Low/Medium – Tools like Microsoft 365 Copilot can draft slides or find images quickly [38] [39]. Faculty input needed to vet accuracy.	Low – Main risk is factual errors or inappropriate examples if AI goes unchecked. IP concerns: using AI-generated content – ensure no copyright issues (or use in-house content only).	**Emerging:** Many faculty started trying ChatGPT for class prep in 2023. By 2024, 22% of faculty were regularly using GenAI tools [40]. High faculty interest, but need best practices to ensure quality. Likely to be mainstream soon.
AI Research Assistant for Faculty (literature review, data analysis help)	Efficiency; Innovation	+Publications per faculty (long-term); -Time to do lit reviews	Academic databases (for summaries); research data sets (for analysis via ML)	Medium – GPT-based summarizers can handle lit review queries. Data analysis assistants (AutoML) require proper data and statistical oversight.	Medium – Risk of missing critical literature if AI has training cutoff or biases. Data analysis AI must be validated (could produce flawed results if used naively).	**Emerging:** Tools like Elicit, Scite, and consensus apps are used to summarize papers. Some faculty use AI to brainstorm or even draft parts of papers (with oversight). Early but growing adoption in 2024.

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk	Maturity
Automated Meeting Notes & Action Items (faculty committees, advising sessions)	Efficiency; Transparency	-Time spent on meeting admin; +Follow-through on actions (clear AI-captured list)	Meeting audio transcripts; participant identification	Low – Many off-the-shelf solutions (Otter.ai, MS Teams with AI notes). Implementation is straightforward.	Low – Ensure confidential meetings are not sent to external AI without consent (privacy). If sensitive, use on-prem or opt-out.	**Proven:** Widely used in corporate, now academia following (e.g. MS Teams with Copilot can summarize meetings [38]). Likely standard practice in near term for non-confidential meetings.
Faculty Teaching Coach (AI analyzes teaching videos or LMS data to suggest improvements)	Quality (Teaching effectiveness)	+Student eval scores; +Learning outcomes in classes where applied	Lecture video recordings; mic audio; student feedback data; LMS usage patterns	High – Requires analyzing complex data (tone, pace, content coverage vs syllabus). Could use AI to transcribe and flag, but actionable coaching needs sophisticated patterns.	Medium – Faculty privacy and autonomy: need consent to analyze their teaching. Risk of over-reliance on AI “rating” teaching quality (should not be used for evaluation without human review).	**Experimental:** Some research efforts to use AI for teaching reflection (e.g. AI detecting if instructor talks too fast, or not enough wait time after questions). Not yet productized at scale. Could be pilots in centers for teaching excellence.

Progression, Alumni & Operations

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk	Maturity
AI Career Advisor (for students & new alumni)	CX; Outcomes (career placement)	+Job placement rate; +Timeliness of job applications; +Student confidence in career plans	Job postings data; alumni career paths; student skills profile; résumé data	Medium – Can use LLM to help with résumé feedback, job matching algorithms, and chatbot for career FAQs. Integration with job boards needed.	Medium – Advice must be accurate/up-to-date (bad career advice can harm futures). Bias: ensure AI doesn't steer certain demographics to lower-paying jobs due to training data bias. Privacy: handle personal career interests data carefully.	**Emerging:** Some universities partner with tools like Handshake' AI or Linked Coach. AI résumé reviewers a common (e VMock). Chatbots fo career Q&A pilot at a fe career centers. Promising results, but still early.
AI Alumni Outreach Personalization	Efficiency; Engagement	+Alumni participation (event attendance, mentorship); +Donation rates (possibly)	Alumni database (interests, past engagement); comms logs; social media mentions	Medium – Similar to marketing: segment alumni and tailor messages. AI can draft personalized emails or find the best channel/ time.	Low – Standard marketing ethics apply. Ensure data usage respects privacy settings alumni have (e.g. if an alum opted out of certain communications).	**Emerging:** Institutions experiment with AI to draft more engaging alumni newsletters (with perso touches like "Congrats o your 5-year grad anniversary Too early to measure donation impact, but engagemen clicks are reportedly u in pilots.

AI Use Case	Value Type	KPI Impact	Data Required	Complexity	Governance Risk	Maturity
AI-Optimized Resource Allocation (timetables, classroom usage, budget)	Efficiency; Cost saving	+Classroom utilization%; -Scheduling conflicts; +Operational cost savings	Facility schedules; course enrollments; utilities usage data; budget data	High – Requires advanced optimization algorithms and integration across registrar, facilities, finance systems. AI can simulate scenarios (digital twin of campus operations).	Medium – If AI recommends cutting certain costs or sections, need human judgment to ensure academic mission isn't compromised. Transparency needed in criteria.	**Emerging:** Some campuses u AI for energ manageme (smart HVAC adjustment saving $$). Academic scheduling is proven in small-scale, full campus optimizatio still in early phases.
ChatGPT (GenAI) Student Assistant (allowed use of GenAI for learning with institutional support)	Learning boost; Integrity (if guided)	+Student productivity on assignments; +Innovation in student work (if using AI for brainstorming); -Incidents of misconduct if guided properly	GenAI tools (like ChatGPT) usage data; guidelines provided; student submissions	Low/ Medium – Allowing use is easy; the work is in providing guidelines and training for *how* to use AI ethically. Possibly integrate AI (like an institutional ChatGPT) with monitoring.	High – Without guidance, GenAI use can lead to plagiarism or fabrication. However, universities are shifting to teaching *with* AI, stressing disclosure and proper use [41] [42] . Need strict academic integrity policies (e.g. requiring citation of AI help [43] [44]).	**Emerging (rapidly):** B late 2023, ~50% of students regularly us GenAI, ever faculty don' [40] . Leadin universities now craft policies to channel this use productively rather than ban it [45] [4 Maturity is mixed: prov that studen *will* use it, b best pedagogica practices ar still being refined.

Legend: CX = Customer (Student) Experience; KPI = Key Performance Indicator.

Each use case above is categorized as Proven, Emerging, or Experimental based on current global adoption and evidence in higher ed: - **Proven** – Demonstrated positive impact in multiple university deployments (high confidence). - **Emerging** – Piloted with some success; becoming more common, but not yet standard practice (medium confidence). - **Experimental** – Limited trials or theoretical proposals; impact uncertain or significant challenges identified (low confidence until further validation).

Reference Architecture – AI-Native University

Design Principles: An AI-native university platform treats AI as a core layer in the enterprise architecture – not a bolt-on to legacy systems. Key principles include a robust data foundation, modular AI services (with centralized model governance), seamless integration with user-facing applications, and a strong governance overlay for ethics and compliance. Below we outline a reference architecture comprising five layers (Identity, Data, Model, Application, Governance) and illustrate how they manifest in three critical domains: the student experience platform, student success & advising, and teaching & assessment.

A. AI-Native Student Experience Platform

Use Case Focus: Personalized student portal and learning environment where AI powers tutoring, content recommendation, and 24/7 assistance (the "AI Student Companion").

- **Identity & Access Layer:** Unified Identity and Access Management (IAM) connects students to all services. Single sign-on ensures the AI assistant knows the user's identity, program, and permissions. For example, integration with Active Directory or a campus SSO means when a student asks the AI "What's my next class?", the assistant can retrieve their timetable securely [47]. Role-based access controls ensure the AI only accesses data the student is authorized to see (e.g. their own grades).

- **Data Plane:** A **Learning Record Lakehouse** aggregates fine-grained interaction data (clickstreams, assessments, forum posts), typically in a cloud data lake. This is paired with an **Event Stream** (e.g. Kafka) capturing real-time events like "student X just missed a deadline". A **Feature Store** serves engineered features to AI models (e.g. cumulative quiz scores, last login frequency), and a **Vector Database** stores embeddings of course content and institutional knowledge (for retrieval-augmented generation in the assistant) [48]. For instance, all course documents (syllabi, slides, transcripts) are indexed in a vector DB so the AI assistant can fetch relevant snippets when a student asks a course content question [35]. Data sources feeding this layer include the LMS, SIS, library, and even IoT (campus card swipes for attendance).

- **Model/AI Plane:** Hosts various AI services:

- **LLM Gateway:** A middleware that routes prompts to appropriate large language models (could be an open-source model hosted on-prem or via API to a vendor). It handles prompt augmentation (injecting retrieved context), and output moderation. For the student assistant, the LLM gateway ensures answers are grounded in approved content – e.g. using retrieval augmented generation (RAG) so that the LLM's output is backed by course materials [35] [6].

- **Domain Models:** beside LLMs, specialized models run here: e.g. a course recommendation model (matrix factorization or neural recommender), a knowledge tracing model to predict mastery level per concept, and an early-alert predictive model for engagement. These might be served via an inference service (SageMaker endpoints, etc.).
- **Agents & Orchestration:** An **AI Orchestrator** (could be based on frameworks like Bedrock Agents or LangChain) enables complex multi-step tasks [34]. For example, when a student asks, "Help me plan my study schedule for finals," an agent might: query their exam schedule from SIS, check upcoming assignment deadlines from LMS, then call the LLM to generate a study plan. Or an agent that, upon a student request, can **book an advising appointment** by interacting with a calendar API after confirming the student's availability.

- **Memory & Personalization:** The platform maintains a **conversation history store** (perhaps a vector store or a simple DB keyed by user) so that interactions with the AI assistant carry context over time (for continuity and learning preferences). For instance, if a student often asks math questions, the system remembers to provide more step-by-step explanations. This "long-term memory" is constrained to prevent sensitive data retention beyond consent.

- **Application Plane:** The student-facing touchpoints: a unified mobile app or portal that embeds AI features natively. The **AI Assistant UI** could appear as a chatbot overlay across systems (LMS, portal, mobile app). Also included are classic apps: LMS, library system, etc., but enhanced to invoke AI services. For example, within the LMS, a "Ask the Tutor" button sends context (course, current page) to the LLM gateway and returns an answer [35]. Or the portal provides a personalized dashboard ("Upcoming tasks you might want to tackle today, and recommended practice quizzes") generated by AI from the student's data. Integration is key: the portal uses APIs to pull data from the Data Plane (via microservices or GraphQL) and to send user actions back as events.

- **Governance Plane:** All AI interactions with students are logged and auditable. **Ethical Guardrails** are implemented at multiple points: e.g. the LLM gateway uses a moderation filter (blocking disallowed content or advising-language that violates policy) [6]. **Consent management** is critical – students opt in to certain AI-driven services (especially those using personal data in new ways). The platform includes an **AI Usage Dashboard** for administrators to monitor queries, response quality, any flagged outputs (e.g. if the assistant gave potentially harmful advice, it's recorded for review). Bias and fairness checks are also run – e.g. monitoring if the course recommender disproportionately suggests certain majors to certain demographics, triggering an alert for review. This plane ties into the institution's broader data governance (FERPA compliance, etc.) and risk management processes (e.g. model risk assessments and documentation).

Example Flow: A student types a question in the mobile app: "I'm struggling with calculus integration – can you help me like a tutor?" The request hits the AI Assistant Service (Application Plane), which authenticates the student (Identity layer) to tailor the help to their course enrollment (Calculus I). The assistant's Orchestrator retrieves the student's recent performance data from the Data Lake (they scored low on last two integration quizzes) and pulls relevant course notes from the Vector DB (Data Plane). It then crafts a prompt with those notes and sends to the LLM (Model Plane). The LLM's response (a step-by-step explanation with a practice problem) is returned to the app. The entire interaction is logged; because this relates to academic work, the governance rules might also send a gentle reminder: "Remember to use this for understanding, not for cheating. Here's your university's AI usage honor code." All of this happens in seconds, delivering a personalized learning moment powered by an orchestrated AI architecture [34] [5].

B. AI-Native Student Success & Advising Platform

Use Case Focus: An integrated system where AI monitors student progress and assists advisors and faculty in intervening to improve success (retention, completion, wellbeing). Think of this as the "AI co-pilot" for student success staff and an early-warning nerve center.

- **Identity & Access:** Here we include not just student identities but also advisor/staff roles. Advisors log into a dashboard that customizes insights to their cohort. Role-based access ensures an advisor sees only their students' data. Students, if they use self-service success tools, see only their own predictive risk or suggestions (and possibly have the option to hide certain sensitive data from AI analysis if they choose, as part of consent preferences).

- **Data Plane:** Emphasis on a **Unified Student Data Hub**. This aggregates SIS data (demographics, grades, credits), LMS data (logins, assignments), advising notes, financial aid status, even swipe data (e.g. dining hall or dorm entry which can flag if a student hasn't been seen on campus for days – a potential concern). Real-time event streams trigger alerts (e.g. "midterm grade below C in 2 classes" event). A **Feature Store** here might maintain up-to-date risk factors for each student (attendance rate, cumulative GPA trend, etc.). A vector DB might store unstructured data embeddings – e.g. past advising transcripts or student reflective essays – that could be queried by an advisor's AI assistant to summarize a student's situation.

- **Model/AI Plane:** Several AI components:

- **Predictive Models:** e.g. a **dropout risk model** and **course failure prediction model** running continuously, outputting scores to the feature store. These are typically ML models (random forest, XGBoost, or neural nets) trained on historical student outcomes [18] [20]. They output risk probabilities along with top contributing factors (for explainability to advisors).
- **Advising Virtual Assistant:** Similar to the student-facing one but tuned for advisors and success coaches. For example, an advisor can query, "Which of my advisees are most at risk this semester and why?" The assistant might use the predictive model outputs plus logic to group by reason (e.g. financial risk vs academic risk) and generate a summary [49]. Or an advisor might ask, "Suggest an action plan for student Alice who is struggling," and the system could retrieve Alice's data and craft a recommendation (e.g. tutoring resources, a meeting invite) based on similar past cases (with guardrails that final decisions are human's).
- **Nudge/Notification Engine:** An AI service that decides which message to send to which student when. This could be rules-based or use reinforcement learning to optimize response rates. E.g., if a student's risk rises, the system might recommend the advisor send a personalized text; if the student hasn't logged into LMS for 7 days, the AI could draft an encouraging email. These drafts would use an LLM for tone ("empathetic and growth-oriented"), but require advisor approval to send (to ensure appropriateness).

- **Conversational AI for Students:** A student success chatbot that checks in on students: "Hey, everything ok? You missed two classes, do you need any help?" This could be automated for certain triggers, escalating to human if the student indicates serious issues. The Model Plane might host an intent-classification model to route wellbeing concerns to counseling versus academic issues to advisors.

- **Application Plane:** For advisors/staff, a **Student Success Dashboard** aggregating alerts and AI insights. It might have a triage view: red/yellow/green statuses with explanations ("High risk: GPA drop and missing financial aid form [50] "). Advisors can click and see an AI-generated summary of that student's "story" (e.g. "Alice, 2nd year, first-gen, missed 30% of classes, financial hold on account – at risk academically and financially"). There is also a workflow for advisors to log outreach, with AI suggesting next steps or resources to send. For students, their interface is through nudges (text, emails, app notifications) and possibly a student-facing portal where they can see, for example, a "Degree Progress AI Advisor" that answers "What if I changed my major?" by running scenarios through the degree audit rules via the AI.

- Additionally, the **Application Plane** might integrate with CRM for scheduling (so the AI can directly schedule a meeting in the advisor's calendar if student responds "I need help"). The platform could also integrate with an LMS early-alert plugin that allows faculty to add context (like "Student X has not submitted last 3 assignments" – which flows to the data layer and triggers AI analysis).

- **Governance Plane:** Because this deals with sensitive student records and life-impacting decisions, governance is paramount.

- **Bias & Fairness:** The predictive models must be regularly audited. For example, the governance layer might include a bias audit report every semester checking model performance across demographics. If any bias is detected (say the model flags low-income students as higher risk mainly due to financial strain – which might be true but requires careful handling), policies might dictate offering additional aid rather than simply labeling them at-risk.
- **Transparency:** Many universities choose to **not** show risk scores directly to students because of stigma. Governance might mandate that risk predictions are *only* seen by advisors, and any intervention must be framed positively (no student is told "You are predicted 80% likely to drop out" – instead they receive supportive outreach).
- **Consent & Privacy:** Students should be informed that their data (LMS, attendance, etc.) is being used to support their success, ideally at enrollment with ability to ask questions. Some data may require opt-in (e.g. using mental health check-in chatbot might be opt-in due to sensitivity). All data and AI usage aligns with FERPA and local privacy laws – likely overseen by the Data Governance Committee and the AI Ethics Board.
- **Human Oversight:** A clear governance rule could be: *No punitive action is ever taken based solely on an AI prediction.* The AI can flag and recommend, but advisors make final decisions and must document their judgment. Also, a process to contest or correct AI-driven records (if a student says "the system thought I was disengaged but I had a family emergency – this needs context").
- **Monitoring:** The system logs all messages sent and outcomes (did student respond, did GPA improve). This allows continuous improvement and checking for unintended consequences (maybe nudges are stressing students out – surveys could be correlated).

Architecture Pattern: The Student Success AI acts like a **"retention radar"**: constantly scanning data for trouble signals, and a **"copilot"**: helping staff prioritize and personalize support [18] [19]. It exemplifies a human-AI hybrid approach: AI for number-crunching and initial outreach at scale, human for relationship and mentorship. The architecture ties together data that was once siloed (academic, financial, engagement) into a cohesive, AI-analysed whole [51] [17].

C. AI-Native Teaching & Assessment Platform

Use Case Focus: Augmenting faculty capabilities in content creation, grading, and maintaining academic integrity, while enabling data-driven teaching. Essentially an "AI Teaching Toolkit" integrated into the learning environment.

- **Identity & Access:** Extends beyond standard IAM to academic roles and course contexts. For example, an "AI Teaching Assistant" (like Jill Watson) operates within the context of a specific course and only users enrolled in that course can interact with it [52]. Faculty have control to turn on/off certain AI features in their courses. Identity also differentiates instructor vs student queries – e.g. a faculty asking the AI to generate a quiz gets full access to content, a student asking likely goes through a different filter (to prevent getting quiz answers).

- **Data Plane:** Central here is a **Learning Content Repository** and **Knowledge Graph** of the curriculum. Course materials (slides, textbooks, past questions) are stored and tagged with concepts. This is fed into a vector index for AI to retrieve context when answering student questions [35]. There's also a **Rubric and Example Repository** for assessments – storing past student submissions, grading rubrics, common feedback points.

- Additionally, a **Telemetry DB** collects how students interact with content (e.g. which pages of e-textbook they spend time on, which quiz questions they get wrong). This feeds learning analytics and allows experimentation (A/B testing data collection).
- **Assessment Records Store:** All assignments, grades, and even intermediate steps (drafts, practice answers) are stored. This is useful for AI models that analyze learning progress or detect integrity issues (like sudden jumps in quality that might indicate AI-assisted cheating).

- If doing automated proctoring or video analysis of presentations, that multimedia data might be processed and stored short-term (with strict deletion policies to protect privacy).

- **Model/AI Plane:** Several AI components geared to teaching/assessment:

- **AI Grading Assistant:** Models like computer vision for scanning handwritten work (like Gradescope does), clustering algorithms to group similar answers [26], and LLMs to draft feedback comments. For programming assignments, a code LLM might evaluate style or efficiency beyond test cases. These models are hosted with secure access to the specific course's data. Importantly, there's a feedback loop: as instructors override or correct AI-proposed grades/feedback, the system learns or at least logs improvements for next time.
- **Content Generators:** E.g. quiz question generators (like D2L's Lumi uses an LLM to make questions from content [14]), slide summarizers, even synthetic labs or case studies. Possibly an **AI Simulation engine** if the course uses scenario-based learning (generate new case variations).
- **Integrity & Authorship Checkers:** Alongside Turnitin's tools, the platform might include an AI that compares a student's writing style across submissions (to flag if a completely different style appears – indicating possible AI use or contract cheating). Or an image forensics AI if students submit images as work. These are separate models specialized in anomaly detection.
- **Teaching Analytics Models:** For example, an algorithm that analyzes LMS data to find which content students struggled with (long time spent but many errors on related quiz) and suggests to the instructor "Module 2 was challenging; consider reviewing it." This could be a combination of

unsupervised clustering (to find difficult content) and simple predictive models linking content engagement to grades.

- **AI TA (Chatbot) for Q&A:** As in Jill Watson's case, an LLM-driven agent that answers student questions on discussion forums [53]. The architecture here (per Jill Watson) uses a curated course knowledge base, a classifier to identify question type, and only answers if confidence is high and content is verified [35] [6]. This AI TA might escalate questions it can't handle to human TAs.

- **Application Plane:** These capabilities surface in the LMS and educator tools:

- Within the LMS, instructors see options like "Autograde submissions" or "Generate feedback draft for all essays." They can review/edit AI-suggested grades and comments before releasing to students (critical for trust). LMS forums show AI responses labeled as such (e.g. "AI TA Alice: ..." with maybe an icon), and give faculty the ability to approve or correct those answers.
- A content authoring interface (could be LMS or separate) where faculty can say "Give me 5 practice questions about topic X" and the AI generates them, pulling from the content repository. The faculty can accept, edit or reject. Instructure's partnership with OpenAI and Khanmigo is heading this direction [54].
- **Student-facing:** When students submit work, they might get immediate AI feedback (if enabled) like "Your essay's introduction is strong, but you might consider adding more examples in the third paragraph." This could be delivered via the LMS assignment interface after submission (as a formative step before final deadline if allowed). Another student-facing element: when taking tests or doing homework, AI proctoring might run in the background (if an online exam, an app might monitor via webcam – though as noted, this is contentious and needs clear communication).
- A "What-if Analyzer" might be available: e.g., a student could ask an AI "What kind of answer is the grader expecting here?" and it could respond with hints (if allowed by instructor). Or "Show me a similar problem" and it fetches one from content library.

- Also, the Teaching & Assessment platform might integrate with the **Student Experience assistant** – for example, the same chatbot that answers general questions can also answer course-specific ones if plugged into this system.

- **Governance Plane:**

- **Academic Integrity Policies:** The system enforces whatever the policy is on AI usage by students. For instance, the platform could have toggles per course: "GenAI Assistance *permitted/forbidden* on Assignment 1" which then informs the AI (if forbidden, the AI TA won't answer a direct question that basically gives away the assignment). It also reminds students of disclosure policies (e.g. if AI is allowed with attribution [43], the system might prompt "Did you consult any AI tools? Remember to cite them.").
- **Model Accuracy & Bias Checks:** Automated grading models must be tested for bias – e.g., does the essay scorer give lower scores to non-native English phrasing? Governance might require periodic calibration (e.g. 10% of exams are double-graded by humans to compare). If discrepancies are high, AI grading is adjusted or suspended.
- **Data Privacy:** Students and faculty should know how their data is used to train/improve AI. Likely, the system is confined to institutional data (not sending student submissions to external AI providers without consent). If external LLMs are used (e.g. via API), governance ensures no personally identifiable info is sent, or a contract assures data not retained by provider [55].

- **Transparency to Faculty:** Faculty need to trust the AI suggestions. The system might provide an "AI decision trail" – e.g. highlight which parts of an answer led the AI to deduct points, or which reference material it used to answer a student's question [6]. This helps faculty audit the AI's pedagogical alignment.
- **Ethical limits:** The governance could set boundaries, e.g. the AI TA should not answer if it detects the question is actually the exact homework problem (to avoid solving assignments for students). This can be implemented via similarity checks. Or in sensitive subjects, AI content generation might be disabled to avoid, say, generating a wrong medical advice in a clinical training course.

In summary, this architecture empowers instructors with AI, but under their control: AI handles grunt work (drafting materials, initial grading, common questions) so faculty can focus on higher-level teaching [56]. It also continuously collects learning data to refine teaching strategies over time. The **model plane** in this context must juggle both generative AI and deterministic models (for grading and analytics), all orchestrated with strong governance oversight to maintain academic standards and trust.

Comparison of LMS-Centric vs SIS-Centric vs AI-Native Models

Many universities today operate in either an **LMS-centric** or **SIS-centric** digital model, and we propose a shift to an **AI-native** model. The table below compares these operating models across key dimensions:

Dimension	**LMS-Centric Model** (Traditional)	**SIS-Centric Model** (Administrative)	**AI-Native Model** (Transformative)
Core System "Brain"	LMS as primary hub for content and student interaction. SIS and others often feed into LMS in a limited way. Decisions revolve around course delivery.	SIS/ERP as the master system (student records, finance, HR). Academics seen through the lens of records and processes.	AI/Data platform at the core – continuously learning from data. The "brain" is a layer of AI services connecting LMS, SIS, CRM, etc., enabling dynamic decisions (personalization, predictions) [49].
Data Integration	Fragmented: LMS holds grades/ activities; separate systems for advising, CRM, etc. Limited analytics combining them (maybe occasional manual exports).	Centralized administratively: SIS integrates with finance and HR well, but learning data might be peripheral. Data used mainly for reporting (credit hours, transcripts).	Unified data lakehouse fusing academic, engagement, and operational data in real-time [57] [58]. Enables holistic analytics (e.g. tying engagement to outcomes). Data is a strategic asset, not just for reporting but for AI-driven actions.

Dimension	**LMS-Centric Model** (Traditional)	**SIS-Centric Model** (Administrative)	**AI-Native Model** (Transformative)
Personalization	One-size-fits-all content delivery. Some adaptive learning if an external tool plugged into LMS. Largely static course experiences.	One-size-fits-all services. SIS treats students as records, not adapting services. Communications are mass or segmented by basic criteria.	Highly personalized: AI tutors adapt content difficulty to each student [4]. Advising AI gives student-specific guidance. Even administrative processes (suggesting personalized course schedules or financial aid nudges) are tailored by AI.
Decision-Making	Human decision-heavy, batch mode: Instructors manually identify struggling students via gradebook; interventions rely on individual effort or static reports.	Policy and rule-driven: decisions (admissions, degree audit) rely on static rules in SIS. Limited predictive insight; reactive interventions (e.g., probation after low GPA happens).	Proactive and autonomous decisioning: AI predicts issues (academic or admin) early and triggers supportive action [18] [19]. Many micro-decisions automated (e.g. routing a question to the best advisor, optimizing course offerings based on forecast demand). Humans focus on oversight and complex cases.
Scalability & Efficiency	Constrained: Faculty and staff workload grows roughly linearly with enrollment. E.g., more students = more grading, more advising appointments, no assistive tech in place.	Efficient in transaction processing (SIS can handle large course registration volumes), but not in individualized support. Past a point, student experience suffers because support doesn't scale individually.	Scales with less friction: AI handles large volumes of repetitive tasks (grading at scale via AI assistance [26], chatbots answering thousands of queries [21]). Human effort scales logarithmically – staff manage by exception rather than every case, allowing growth without equivalent cost explosion.
Innovation Cycle	Slow: LMS releases features occasionally; adding new tools requires lengthy integration and faculty training. The tech is often behind consumer tech.	Slow: SIS-centric environments prioritize stability over new features. Upgrades to core SIS are infrequent and disruptive. Innovation mostly on periphery (small pilot tools).	Fast and experimental: The AI-native model encourages continuous experimentation (A/B testing of interventions, new AI features trial). Modular architecture allows plugging in improved models or tools quickly. There is an ethos of ongoing optimization (like tech companies).

Dimension	**LMS-Centric Model** (Traditional)	**SIS-Centric Model** (Administrative)	**AI-Native Model** (Transformative)
Student Experience	Course-focused and fragmented: students jump between LMS, portals, email, asking staff in person for some tasks. Many friction points (e.g., not knowing where to get answer).	Process-focused and bureaucratic: students see the university as a set of offices/processes to navigate (admissions, registrar, etc.). Experience depends on individual initiative to navigate silos.	Seamless and concierge-like: a single AI-enhanced interface can handle a wide range of queries ("one-stop"). Student gets immediate support or is guided to the right resource by AI. Administrative and learning support blend into one personalized support system. The experience is more "on-demand" and Amazon-like (in responsiveness), which students increasingly expect [59] [60].
Faculty & Staff Role	Faculty: primarily content delivery and grading. Staff: transaction processing (enroll students, schedule, etc.) plus reactive support. Some skepticism of tech because it hasn't clearly reduced their burden.	Faculty: providers of grades that feed SIS. Staff: guardians of process and records. Roles siloed (advisors advise, IT runs systems, etc.), focus on compliance and policy enforcement.	Faculty: freed from many low-level tasks (AI helps with grading, materials), they can spend more time on mentoring and research. They also use AI insights to refine teaching (learning analytics). Staff: elevated to "strategic facilitators" – e.g. advisors manage by exception with AI highlighting who truly needs help. New roles emerge (AI trainer, prompt curator, ethics reviewer). More cross-functional teams (IT, academic, student affairs collaborating on AI initiatives).

Dimension	LMS-Centric Model (Traditional)	SIS-Centric Model (Administrative)	AI-Native Model (Transformative)
Governance & Risk	Often ad-hoc: Each system has its own data use and security rules. Little centralized AI or data governance because usage was limited. Risk of data privacy issues if systems aren't well integrated (shadow databases, etc.).	Policy-driven: Strong governance on data that goes into SIS (FERPA compliance top of mind). But less so on learning data outside SIS. AI not a focus, so no specific governance for it (except maybe plagiarism tools).	Holistic AI governance structure in place. Clear policies on ethical AI use, bias audits, transparency to students [61] [62]. A governance committee oversees all AI projects. Compliance extends beyond FERPA to cover AI ethics, model validation, security of new data types (e.g. voice transcripts). High awareness of risks like algorithmic bias and proactive measures to mitigate them.

In summary, **LMS-centric** models center on delivering content but often lack deep data insight or personalization; **SIS-centric** models ensure administrative control and consistency but can be impersonal and slow to adapt. The **AI-native model** combines the best of both (rich learning interactions + robust data backbone) and adds a layer of intelligence that makes the university more responsive, personalized, and scalable. This requires re-thinking systems not as isolated pieces but as an integrated intelligent ecosystem.

Operating Model for an AI-Native University

Transforming into an AI-native model isn't just about technology – it demands a new operating model, including organizational structures, roles, and processes. We propose establishing a dedicated **University AI Office (UAIO)** and associated governance structures, with clear RACI (Responsible, Accountable, Consulted, Informed) assignments for key processes.

University AI Office (UAIO): A central team driving AI strategy, adoption, and oversight. This office works cross-functionally (IT, academic affairs, student services, IRB, etc.). Key roles within or connected to this office might include:

- **Chief AI Officer / AI Program Director** – *Accountable* for overall AI transformation strategy and outcomes. Chairs the AI governance committee. Ensures alignment with institutional goals. (RACI: A for AI strategy, A for AI policy compliance)
- **AI Solution Architects** – *Responsible* for designing the AI platform and integrations (technical blueprint, reference architecture enforcement). Work under CIO/CTO but liaise with academic units. (RACI: R for architecture design, C for selecting AI vendors/partners)
- **Data Scientists / ML Engineers** – *Responsible* for developing and tuning models (predictive analytics, LLM fine-tuning, etc.). Some embedded in IR (institutional research) or IT, but coordinated via UAIO for standards. (RACI: R for model development, R for validating model performance; C with faculty if academic data is involved)

- **AI Ethicist / Policy Lead** – *Responsible* for governance and ethical oversight. Develops guidelines for AI use, runs bias audits, manages training on AI ethics for faculty/staff [61] [62]. (RACI: R for bias audits, A for ethical approval of AI projects in academics; I faculty on findings)
- **Product Managers (AI Initiatives)** – *Responsible* for specific AI projects (e.g. AI tutor rollout, AI advising system). They coordinate between users (students/faculty) and tech teams, manage timelines, training, feedback loops. (RACI: R for project execution, C with end-user departments, I to AI Officer on progress)
- **Integration Specialists / Data Engineers** – *Responsible* for connecting legacy systems (LMS, SIS, CRM) to the AI platform and maintaining data pipelines. (RACI: R for data integration quality, C with system owners like Registrar (SIS) or LMS admin)
- **Change Management & Training Lead** – *Responsible* for the human side: plans faculty/staff development around AI, documents best practices, collects feedback. Possibly from Center for Teaching and Learning working with UAIO. (RACI: R for training programs, C with faculty senate or union if needed)
- **Cybersecurity & Privacy Officer** – *Consulted/Responsible* to ensure AI systems meet security standards, data privacy laws (e.g. Vietnam's PDPL, GDPR if international students). (RACI: C on all AI implementations, A for security clearance to go-live)

Governance Bodies and Processes:

- **AI Governance Committee:** A cross-functional committee that includes the AI Officer, CIO, faculty representatives, student representative, legal counsel, and possibly an external ethics advisor. *Accountable* for approving high-level AI policies and major deployments (like using AI in admissions). Meets regularly to review AI use cases, risk reports, and compliance with principles (fairness, transparency, etc.) [61]. (RACI: A for policies; R = AI ethicist prepares materials; C = faculty/students consulted on academic impacts)
- **Data Governance Subcommittee:** Focused on data quality, sharing, and consent issues for AI. Ensures proper data stewardship for training AI models. (Likely overlaps with existing data governance board but with AI specifics). (RACI: A = Chief Data Officer or CIO; R = Data engineers ensure anonymization as needed; I = legal and ethics)
- **Academic Senate / Curriculum Committee role:** They need to be *Consulted* and *Informed* for AI that impacts teaching and learning. For example, if AI auto-grading is introduced, academic bodies set parameters (like "AI feedback can count for participation but final grades by instructor"). The faculty governance ensures pedagogical soundness. (RACI: C for academic policy on AI in classrooms; A = Provost on final decisions incorporating committee input)
- **IT Change Management** (existing ITIL processes): incorporate AI changes. E.g., any new AI system goes through review similar to a new SIS module, including security and user acceptance testing.

Key Processes & RACI Examples:

- *Identify AI Use Cases:* UAIO Product Managers work with departments to identify opportunities. **Responsible:** Product Manager; **Consulted:** Department heads, faculty, students (end-users); **Accountable:** Chief AI Officer to approve pursuing it; **Informed:** CIO, Provost.
- *Develop/Procure AI Solution:* **Responsible:** AI Solution Architect & Data Scientist; **Accountable:** CIO/CTO for tech decisions (with AI Officer); **Consulted:** End-user reps (to ensure requirements met), Purchasing (if vendor involved); **Informed:** AI Governance Committee (if high-risk).
- *Bias & Impact Assessment (for a new model):* **Responsible:** AI Ethicist and Data Scientist to produce bias report; **Accountable:** AI Governance Committee to approve proceeding; **Consulted:** Legal (for

compliance), Affected department (e.g. Admissions if model used there); **Informed:** University leadership and perhaps public transparency if required.

- *AI System Deployment:* **Responsible:** Integration Specialist & IT ops; **Accountable:** CIO; **Consulted:** Security Officer, Data Privacy Officer; **Informed:** All end-user stakeholders (with a clear briefing on how to use it, any policy changes).
- *Monitoring & Incident Response:* e.g., AI gave a problematic output or there's a potential breach. **Responsible:** UAIO team to investigate; **Accountable:** Chief AI Officer to coordinate response (with CIO for technical, and Ethics for handling user impact); **Consulted:** Affected department and Comms (if messaging needed); **Informed:** Governance Committee, possibly regulators depending on severity.
- *Continuous Improvement:* Quarterly review of AI outcomes (e.g., is student retention up? Are faculty using the tools?). **Responsible:** Data Analyst in UAIO to compile metrics; **Accountable:** AI Officer to adjust strategy; **Consulted:** Provost, student success VP, etc.; **Informed:** Board or President in strategic updates.

Cultural and Skill Considerations: The operating model also implies upskilling many roles: - Faculty development programs on using AI in pedagogy (led by CTL with UAIO input) – so faculty become partners in refining AI (e.g., co-creating content with AI) rather than feeling displaced [63]. - Staff training on interpreting AI outputs (e.g., advisors reading risk scores appropriately – not as deterministic labels but cues). - Possibly new **AI Student Fellows or Interns** – students involved in testing and providing feedback (both to get student perspective and to train future talent).

Finally, ensure **RACI clarity** in documentation for each major AI system. For instance: - The AI Advising Chatbot: **A:** Director of Advising, **R:** UAIO Data Scientist & Advising lead, **C:** Advisors, **I:** students (via announcement). - The AI Grading tool: **A:** Vice Provost or Dean of Faculty, **R:** CTL lead & IT, **C:** Faculty user group, **I:** all instructors before semester start.

This structure prevents "AI initiatives" from being siloed experiments and instead makes it an institutional capability with oversight and continuous alignment to goals.

36-Month AI Transformation Roadmap

A phased roadmap (0–12 months, 12–24 months, 24–36 months) will guide the transition to an AI-native operating model. Each phase covers platform development, data foundation, AI use case rollout, organizational and skills changes, and governance maturation:

0–12 Months: Foundation and Quick Wins

- **Strategy & Buy-in:** Form the University AI Office (UAIO) and governance committee in first 3 months. Conduct executive education sessions for Board, CEO, leadership on AI-native vision to ensure top-down support. Develop a clear definition of "AI-native EdTech" as per our focus (personalization, autonomous decisioning, etc. – moving everyone beyond "just a chatbot" thinking).
- **Current State Audit:** Inventory all existing systems (ERP, SIS, LMS, CRM, etc.) and digital data silos. Identify data quality issues and integration gaps. Also survey faculty and student readiness (attitudes, skill gaps) to tailor change management.
- **Data Platform Launch:** Stand up a cloud-based data lakehouse for the university (if not already) [57]. Focus on centralizing student data: ingest SIS (student records), LMS logs, CRM (recruitment

info) into the lake in raw form. Implement an initial feature store for simple use cases (like a flag for "attendance < 50%").

- **AI Infrastructure:** Establish basic AI infrastructure in sandbox mode. E.g., subscribe to an LLM service (OpenAI, Azure OpenAI, or local like FPT AI if VN) for pilot use with proper data agreements [55]. Acquire or repurpose a small GPU server for any in-house model experiments. Ensure IT is up to speed on supporting AI workloads (containerization, model ops).
- **Quick-Win Use Cases:** Target 2-3 low-risk, high-value pilots:
- *AI Admissions Chatbot:* Implement a chatbot for answering applicant FAQs (could use a vendor like Mainstay or an in-house using an LLM with admissions Q&A fed in). This is relatively straightforward and addresses a pain point (inquiries overload). Goal: reduce response time and summer melt (expectation: replicate GSU results of >20% melt drop [22]). Timeframe: by month 6, live for next admission cycle.
- *AI Writing Feedback for Students:* Enable a tool like Grammarly or an LLM-based feedback system in the LMS for a few writing-intensive courses or a writing center. Goal: help students improve drafts before final submission, and gauge academic integrity implications in a controlled way (with honor code guidelines). This is a "win" in student eyes (they get help) and helps us gather data on how AI is used. Rollout in pilot by month 9.
- *Predictive Early Alerts (limited):* Take an existing analytics model (or simple rules) to identify first-year students struggling (e.g., missed 2 assignments + exam score <50%). Use existing data to start flagging and have advisors intervene manually. This doesn't require complex new tech if we use current LMS and SIS data with maybe Excel or basic BI. It builds momentum by showing data-driven support improving outcomes.
- **Skills & Culture:** Begin faculty and staff consultations. Identify "AI champions" among faculty – perhaps those already experimenting with ChatGPT in teaching. Engage them in pilot design (for instance, one champion helps integrate the AI writing feedback in her class). Provide initial workshops: "Introduction to AI in Teaching" and "Data-informed advising 101." The tone is experimental and supportive.
- **Governance Setup:** Develop and publish an initial AI Ethics & Usage Guideline. For example, explicitly allow or disallow certain uses of generative AI in coursework (some universities did this in 2023 to set interim rules [64] [45]). Establish policy on data consent – e.g., update student handbook or consent forms to cover use of their data in AI systems (with anonymization where possible). Also, ensure compliance with Vietnam's data laws (the new PDP decree) – likely requiring opt-in for certain personal data uses.
- **Tech Decisions – Build vs Buy vs Partner:** By month 6, decide on which platform components to build internally vs buy. For example, decide if the data lakehouse will be built on AWS (partner with AWS, maybe using their AI Edu framework [65]) or Azure, etc.; decide if LMS AI features will be used via existing vendor (Canvas AI) or if a custom AI layer will be built on top. Early partnership discussions (e.g., invite Microsoft or AWS to do an AI workshop for your team – many are eager to help education now).
- **Milestone at 12 months:** Have demonstrated at least one tangible AI service to students or faculty with positive feedback (e.g., "the new admissions chatbot answered 5,000 questions this cycle and our enrollment yield went up 3% [50]"). Also, have the core data environment in place to support next phases (even if not perfect, the lakehouse exists and key data streams are flowing). Report progress to board with some early metrics and anecdotes.

12–24 Months: Expand and Integrate

- **Scaling Successful Pilots:** Take pilots that showed value in Year 1 and scale them:

- Admissions chatbot -> extend to cover current student FAQs (merge into a general university FAQ bot accessible in student portal). Possibly bilingual capability (English/Vietnamese) to serve international and local students. Use early interactions to improve its knowledge base.
- Early Alerts -> Implement a full-featured predictive analytics tool (could buy Civitas or similar, or build using your data science team). By month 18, have a predictive model for year-one retention using at least 2 years of historical data for training. Integrate its output into advisor workflow (e.g., via advising CRM or a simple dashboard).
- AI writing feedback -> If positive, expand to more courses. Also consider adding AI for other subjects (like a math homework helper pilot using an adaptive tool).
- **New Use Cases Rollout:**
- *AI Advising Assistant (Phase 1):* Develop a knowledge base of advising Q&A (academic policies, program requirements) and pilot an advisor-facing GPT that can draft answers or prep materials for advising sessions. For example, before an advising meeting, the AI summarizes the student's situation and suggests topics. Test with a small group of advisors and refine.
- *Autograding in Large Classes:* Choose a high-enrollment course (perhaps an intro programming or large general ed) and introduce AI-assisted grading (Gradescope or similar) to handle exams or assignments. Monitor grading consistency and faculty satisfaction. If results in significant time savings (e.g., grading time cut by 40%), plan to extend to other courses.
- *Personalized Learning Path in a Course:* Partner with a content provider or use an in-house team to create an AI-driven module in one course (e.g., an adaptive learning module for an online math or language course). The system should adjust to each student's level. Evaluate impact on outcomes vs a control group.
- **Data & Platform Maturation:** By mid-point of this phase, implement the **feature store** and real-time pipelines more robustly. For example, integrate streaming of LMS data (daily or hourly ingestion instead of end-of-term). Deploy a **vector database** and start indexing university knowledge (policies, course syllabi, library FAQs) – this will power more advanced Q&A and tutoring use cases. Possibly adopt an enterprise search solution that is AI-friendly for internal knowledge.
- **Architecture & Integration:** Start refactoring or upgrading legacy systems to be more API-driven for AI integration. For instance, ensure SIS can both feed data to the lake and consume outputs (like writing back an "AI advisor note" or a flag). Work on single sign-on across new AI apps and existing systems to ensure seamless user experience.
- **Buy/Partner where needed:** If building in-house proved slow in year 1, consider buying for year 2. For instance, if early-alert modeling was challenging, possibly partner with an established vendor (Civitas, etc.) to accelerate (with an eye to customizing for local context – maybe a local EdTech in ASEAN has a solution to partner on). Similarly, evaluate LMS vendor offerings: Canvas and others will by now have more AI features – decide if you enable those or continue developing your independent ones.
- For Vietnamese language support, consider partnering with local AI firms (e.g., FPT, Viettel, or AI centers at Viet universities) to develop Vietnamese NLP capabilities for chatbots or content analysis – by 18-24 months you'll know how big a gap language is.
- **Skills & Org:**
- Expand UAIO if needed (maybe initial team was small). Hire or train more data scientists or analysts as the number of models grows. Possibly create an "AI Student Interns" program – get CS students or others to contribute, which also builds talent and buy-in.
- Conduct regular training for end-users: by now some faculty may be using AI in class (e.g., allowing ChatGPT for assignments). The CTL can host sessions like "Teaching in the era of GenAI" and share internal case studies. Provide advisors training on interpreting the predictive dashboard (e.g., a workshop on "Using data to enhance advising").

- Address change resistance actively: Year 2 is often where initial excitement fades and day-to-day reality sets in. Some faculty might complain about AI mistakes or fear job impact. Use the successes and data to show positive outcomes [63], and keep open channels for concerns. Possibly establish an AI advisory council that includes skeptics to voice and address issues (better to have them inside giving feedback).
- **Governance & Policy:**
- Based on year 1 experience, refine AI usage policies. For example, if many students started using AI in assignments, update academic integrity policies university-wide: define what's allowed, how to cite AI, consequences of misuse (taking cues from global best practices [64] [46]). Communicate clearly at semester beginnings.
- Implement formal **AI ethics review** process: any new AI tool being deployed (especially those affecting student grading or opportunities) should go through a bias and risk assessment. By month 24, have done at least one thorough audit (e.g., check the retention model for bias).
- Data governance: ensure any data-sharing with vendors is compliant; e.g. if using cloud, ensure contracts cover data security. Possibly create a role of "AI compliance officer" under the privacy office by this time, as AI use is now more pervasive.
- **Key Deliverables by end of 24 months:**
- Integrated data platform with at least 70% of needed data flowing (including less commonly used sources like co-curricular involvement perhaps).
- At least 5-10 AI use cases in production (mix of student-facing and staff-facing). E.g., Chatbot, early alerts, AI tutor in some courses, automated grading in some courses, personalization in at least one area, etc.
- Metrics to show progress: e.g., first-year retention improved by X% (maybe small initially, but directionally positive) [19]; average time to respond to student inquiries down from hours to seconds for common questions; faculty grading time saved equating to Y hours freed for research; student satisfaction on advising improved in survey.
- Perhaps a recognition: be it internal (leaders praising teams) or external (media noting the university's innovation – which can attract students and partnerships). HolonIQ or other EdTech org might list the university as a case study if things go well.

24–36 Months: Full Transformation and Optimization

- **Enterprise-wide AI Integration:** By year 3, AI is woven into most core processes:
- *Teaching:* Aim for at least 30-50% of courses to be using some form of AI. Not necessarily all having AI TAs, but could be AI-generated practice quizzes, or at least faculty using AI analytics to adjust teaching. Particularly, scale up the AI tutoring/personalization to all large foundational courses (since those benefit most from scaling and consistency).
- *Student Support:* The AI concierge/chatbot should handle the majority of Tier-1 questions across departments (IT help, registrar Qs, financial aid FAQs, etc.). Integrate it with voice interfaces if possible (for accessibility) and ensure it covers Vietnamese and English. Possibly extend to a physical presence (kiosk or robot) on campus for common queries.
- *Advising:* The predictive model is now refined with a few years of data, improving accuracy. It's integrated with interventions (like a student at risk automatically gets enrolled into a supplemental coaching program or is offered a specific support package, with advisor oversight). Possibly implement an AI-driven *degree pathway optimizer* that suggests an optimal set of courses for on-time graduation given student's history – and advisors use this in their planning sessions.

- *Operations:* Start using AI in back-office too: e.g., AI to optimize course scheduling (reduce clashes and underfilled sections), AI for finance (maybe predictive forecasting of enrollment to inform budgeting), and even HR (AI to scan resumes for faculty hiring or to automate repetitive tasks).
- *Continuous Experimentation:* Institutionalize A/B testing and rapid iteration. For example, test two versions of a nudge to see which improves class attendance more [66] [67]. Run pilots of new AI tech every semester in small areas, measure results, then scale or drop. This requires cultural maturity: acceptance that not everything will work, but it's a learning approach (like how internet companies operate).
- **Full AI-Native Architecture Deployment:** The reference architecture described is fully in place. By end of year 3, the university should have:
- A robust **AI platform** (likely on cloud) that is scalable, secure, and cost-controlled. Multiple AI models running (both batch predictions and real-time LLM services) with a monitoring dashboard (tracking latency, usage, cost). Possibly leverage a model-management solution to orchestrate model updates and A/B tests.
- Identity & data integration such that a student or faculty can move across systems and the AI context follows. E.g., the AI assistant in LMS knows the same student profile as the one in advising portal.
- Governance tools: e.g., automated logging and audit for any AI decision that affects grades or admissions, with easy retrieval if needed for appeals or review.
- **Costs**: By now, pay attention to AI compute costs – incorporate cost-control measures (the policy might route heavy queries to cheaper models unless high stakes). Possibly develop some in-house models for specific tasks to reduce reliance on expensive API calls (if that was an issue).
- **Localize and Optimize for Vietnam/ASEAN:**
- Over 36 months, gather data to fine-tune for local context. If using English-centric models initially, by year 3 consider fine-tuning an open LLM on Vietnamese academic data to have an internal model that can handle local language queries better – reducing reliance on outside services and handling cultural context (like local history or literature questions).
- Adjust to local pedagogy: e.g., if VN classes typically use more teacher-centered approaches, maybe the AI tutor is used to facilitate more active learning gradually (with training to students on how to use it effectively, since that might be new). Ensure content recommendations include local content (not just Western OER).
- Compliance with local regulations: Vietnam is likely developing AI regulations – ensure by year 3 you're aligned or even ahead with ethical compliance (could become a national leader example).
- **Partnerships:** By this stage, consider deeper partnerships: e.g., with other universities in ASEAN to share non-competitive data or co-develop AI for education in local languages; with EdTech startups (maybe even spin out successful internal tools as products). Perhaps join international consortia on AI in higher ed to keep at frontier.
- **Organization:**
- The UAIO might evolve into a permanent department (if not already) or be embedded in existing structures (like under the CIO or Provost). Ensure succession planning – initial champions from 3 years ago might rotate, so institutionalize knowledge (documentation, perhaps an "AI Center of Excellence" that remains).
- Many staff roles will evolve: advisors become more data-savvy success coaches, IT staff include ML engineers. Possibly introduce a **Chief Data Officer** if not had, to steward the data which fuels AI.
- Faculty evaluation and workload norms might adjust: if AI grading saves time, perhaps faculty are expected to spend that time on more research or more student mentorship. Open discussions on these to update policies fairly.
- **Culture & Community:**

- By the end of 36 months, aim for widespread acceptance that AI is a normal part of the university's operations (much like LMS became standard over the 2000s). Continue addressing concerns – e.g., academic integrity with AI: shift focus to designing assignments AI can't easily solve and teaching students how to use AI as a tool (not a cheating shortcut) [25]. Essentially, the narrative should be "AI is enhancing, not replacing human education" with numerous internal examples to point to.
- Highlight success stories: the student who was going to drop out but was saved by an AI-prompted intervention, or the faculty who published more research because AI took some teaching load, or overall improved student satisfaction ratings.
- **Evaluation and Adjust Course:** End of year 3 is time for a comprehensive review. Did we meet the ambition to transform the operating model? Evaluate against the KPIs set (retention, graduation, employment, student satisfaction, faculty productivity, cost efficiency, etc.). Identify areas that did not achieve expected gains and investigate why (was the tech not mature? adoption issues? data issues?). Also identify new opportunities that emerged (maybe student demand for AI-related academic programs increased due to our profile – leading to new programs or micro-credentials in AI).
- **Beyond 36 Months:** Develop the next strategic plan, possibly moving from transformation phase to continuous innovation phase. Likely by this time, the university could aim to be a regional leader in AI-driven education, opening opportunities (attracting both students and partnerships). Plan how to keep the momentum: perhaps aim for an AI-driven lifelong learning platform for alumni next, or more advanced use of AI in research (beyond education domain) as a spillover.

Summary of 36-mo goals: By month 36, the university operates as a cohesive intelligent campus. Every student has an AI-enhanced personal learning journey. Every faculty and staff member has AI co-pilots for routine tasks. Key metrics like retention, time-to-graduation, and student engagement have improved noticeably. The data shows fewer students "falling through the cracks" because the system catches them and the community (with AI help) supports them. Cost-wise, efficiency gains from automation are starting to show (perhaps staff can be redeployed to new initiatives instead of repetitive tasks). Risks are under control through strong governance – no major scandals or breaches, demonstrating that the AI-native approach can be done responsibly [68] [25]. This sets the foundation for continued innovation beyond 3 years, with the institution capable of adapting to whatever new AI advances come next.

Executive Summary

Context & Ambition: In an era where AI is reshaping higher education globally, our multi-campus university in Vietnam aims to transform into an **AI-native university** within 36 months. *AI-native EdTech* means AI is at the core of our operating model – driving personalized learning, autonomous decisions, and optimized processes – not just an add-on tool [69]. This strategy aligns with our mission to improve student success, faculty productivity, and institutional agility, positioning us as an innovative leader in ASEAN. Below we summarize key insights, strategic options, and recommendations, along with major risks and first investments.

Ten Key Insights

1. **Global Momentum:** Universities worldwide are adopting AI for tutoring, advising, grading, and more. Over half of students already use genAI tools, outpacing faculty adoption [40]. Leading institutions (Harvard, Stanford, etc.) now embrace AI and focus on guiding its use rather than banning [45] [46]. The window is ripe for us to leapfrog by learning from these early movers.

2. **AI-Native Platforms Exist:** A landscape of AI-powered EdTech solutions is available. AI-native tutoring systems (e.g. Cognii, Squirrel AI) can personalize learning with proven efficacy (up to 2× improvement in mastery in studies [4]). AI-enabled legacy systems (Canvas, Blackboard) are rapidly integrating AI assistants for instructors [8] [70]. We can mix best-of-breed vendors with in-house innovations (like Georgia Tech's Jill Watson TA improved online class performance [5]) to build our ecosystem.
3. **Use-Cases Span the Student Lifecycle:** We identified ~30 AI use cases from recruitment through alumni. Some are **proven** (chatbots reducing summer melt by >20% [22]; automated grading saving faculty time; predictive analytics boosting retention [20]). Others are **emerging** (AI advisors, AI career coaching) with promising pilots. This portfolio approach ensures impact at each stage (recruitment, learning, success, operations).
4. **Reference Architecture:** A robust architecture underpins AI-native operations – *Identity*, *Data*, *Model*, *Application*, *Governance* layers working in concert. Key design: a unified data lakehouse and feature store to fuel AI models; an AI orchestration layer (with LLMs + knowledge bases) to deliver context-aware services [48]; seamless integration into user-facing systems (LMS, portal, CRM); and an ethics/governance plane enforcing privacy and fairness. This architecture is scalable and modular, allowing rapid deployment of new AI capabilities.
5. **Personalization & Autonomy as Differentiators:** The AI-native model enables mass personalization – every student gets a tailored learning path or support journey [71] – and autonomous agents to handle routine decisions (scheduling, FAQs) in real-time. This can markedly improve student experience (e.g., 24/7 AI support so no question goes unanswered [21]) and efficiency (staff focus on complex cases, AI handles repetitive tasks).
6. **Data is Fundamental:** We must invest in data foundation early. AI outcomes are only as good as the data fed. That means breaking silos (ERP, LMS, SIS, etc.), cleaning and linking data, and even collecting new data (e.g., learning interaction telemetry, which currently is uneven). Governance around data consent and sharing is critical – we need student trust that their data is used to help them (and never to harm or surveil inappropriately) [25].
7. **Ethics & Integrity are Make-or-Break:** AI in education raises valid concerns – bias in models could reinforce inequities, or unchecked genAI use could undermine academic integrity. Best practices from peers: involve faculty in redesigning assessments to be "AI-proof" or AI-leveraging, use AI detectors cautiously (Cornell advises against relying solely on them [24]), require disclosure of AI use by students [43], and maintain a human in the loop for high-stakes decisions (admissions, grading). A strong governance framework (bias audits, transparency to users, ethical guidelines) is non-negotiable to navigate these issues.
8. **Faculty & Staff Enablement:** For successful adoption, faculty and staff must be brought along. Initial resistance can be overcome with training and by demonstrating that AI can *augment* rather than replace their roles [63]. Example: showing faculty that AI can handle mundane grading so they can focus on interactive teaching – at Georgia Tech, AI TA freed human TAs for deeper student engagement [56]. We need an "AI literacy" push internally, similar to how top universities are now requiring AI fluency for students and faculty [72] [73].
9. **Proven ROI Potential:** While AI is new, early results indicate tangible benefits – improved retention (potentially 5–10% increase in year 1 retention with proactive analytics [18]), higher enrollment yield (our chatbot could increase yield a few percentage points like GSU [22]), efficiency gains (AI assistants could cut administrative workloads by 20–30%). These translate to financial gains or savings (retaining students boosts tuition revenue; automation can contain staffing costs as we scale). Long-term, an AI-enhanced reputation can attract more students and partnerships.
10. **Local Adaptation:** We must adapt global solutions to Vietnam/ASEAN context. Language localization is key – ensuring AI works in Vietnamese (for both content and user interface). Also pedagogy: our

students may need guidance to use self-directed AI tools if coming from more traditional backgrounds. Regulatory environment is evolving (PDPL privacy law) – we should aim to exceed compliance, setting a benchmark in responsible AI. Not everything needs to be built from scratch here; we can localize existing platforms (e.g., feed Vietnamese curriculum data to an AI tutor) and focus in-house efforts where unique (e.g., Vietnamese language NLP, local student behavior patterns).

Three Strategic Paths

We outline three strategic options varying in aggressiveness, each with pros and cons:

- **Conservative Path (“Guarded Evolution”):** Focus on a few low-risk AI improvements to existing systems, done slowly and with tight control. For example, enable some LMS AI features (auto-grading, simple analytics) and use a vetted chatbot for FAQs, but avoid touching high-stakes areas like grading or admissions with AI initially. This path minimizes risk and faculty pushback since it doesn’t radically change processes. Budget impact is moderate. However, it may fall short of “transformative” – we’d likely make incremental gains but not a step-change in personalization or efficiency. Risk: being too slow while peers leap ahead, and failing to meet the Board’s ambition of an AI-native model.
- **Balanced Path (“Steady Transformation”):** Pursue a moderate number of AI initiatives across different domains with careful monitoring. Prioritize proven use cases (early alerts, chatbots, auto-grading in large classes) for first 1–2 years, then gradually add more innovative ones (AI advisors, adaptive learning) once confidence is built. Invest in data platform and governance from the start to enable scaling. This approach spreads benefits broadly (academics, admin, student services all see improvements) and manages risk via pilots and evaluation phases. It requires solid project management but is feasible within 36 months with current team + some hiring. Likely this path yields strong improvements in student success and satisfaction, though perhaps not the maximum theoretical because we intentionally pace ourselves. It balances innovation with caution.
- **Bold Path (“Leapfrog to AI-Native”):** Aggressively implement AI everywhere possible, aiming to be a showcase “university 4.0”. That means fast-tracking the full architecture, large-scale use of generative AI in curriculum, even considering radical ideas like AI-driven admissions screening (with human oversight) or replacing certain legacy systems entirely with AI-powered platforms. This path could catapult us ahead, potentially achieving dramatic efficiency and a highly personalized student experience quickly. However, it carries high risk: it demands significant investment (both $$ and talent), cultural resistance could be high if change is too sudden, and there’s little room for error if something fails (e.g., a model error could have wide impact). It also assumes a tolerance for experimental failures as we push the envelope. If executed well, we become a leader regionally; if poorly, we could face backlash or system disruptions.

Recommendation: Adopt the **Balanced Path**. It offers meaningful transformation in 3 years without courting undue risk. We will aggressively build the data/AI backbone and implement key AI services, but we’ll do so pragmatically: piloting, evaluating, and scaling what works. This approach is likely to gain broad stakeholder buy-in – we can celebrate quick wins to build momentum, while addressing concerns on a manageable scale. It also keeps options open: if some bold idea proves safe in pilot, we can incorporate it; if some risk emerges, we can adjust before it’s widespread. The Balanced Path ensures we meet the Board’s ambition of true AI-native operation by year 3, evidenced by measurable improvements, in a sustainable way.

Top Risks and Mitigations

Even with a balanced strategy, we must actively manage risks:

1. **Data Privacy Breach or Misuse:** Using more data and AI means higher stakes for data security. *Mitigation:* Strict access controls, encryption, and compliance audits on the data lake and AI systems. Implement privacy-by-design (e.g., anonymize data for model training where possible). Provide opt-outs for students for highly sensitive data usage. Regularly review data retention and purge policies – don't store what we don't need. Engage legal early to ensure alignment with laws (FERPA, PDPL) – e.g., ensure our student chatbot vendor contract specifies no data mining of our info [55].
2. **Algorithmic Bias/Unfair Outcomes:** AI models might inadvertently favor or disfavor groups (e.g., a retention model flagging disproportionate number of lower-income or minority students). *Mitigation:* Establish bias testing as part of model development. Include diverse data in training sets. Involve campus diversity/equity officers or social science faculty in reviewing AI decisions. If a model is found biased, adjust inputs or thresholds, or even decide certain predictions (like in admissions) won't be used to avoid bias. Also ensure AI augmenting human decisions doesn't mask bias ("human-in-the-loop" should have bias training too).
3. **Academic Integrity Erosion:** Students might misuse AI (cheating via ChatGPT, etc.), potentially undermining learning and our reputation. *Mitigation:* This is as much cultural as technical. Develop clear honor code extensions for AI use (allowed vs prohibited) and educate students on them [43] [44]. Design assessments that require personal input or oral defense so that AI-generation alone won't suffice. Use AI detection as a secondary measure but not sole proof (given detectors' unreliability [24]). Provide faculty support to redesign curriculum for an AI-pervasive world (maybe embrace AI for rough drafts but then have in-class writing, etc.). Essentially, turn AI from an enemy to a teaching tool: e.g., require students to critique an AI's answer rather than just asking them to produce an answer.
4. **Faculty/Staff Resistance and Skill Gap:** If faculty feel AI is imposed or threatening, they may not adopt it, dooming projects (e.g., if they refuse to trust AI grading, it won't save time). *Mitigation:* Extensive change management – involve faculty in pilot design, address their concerns (maybe give opt-in choices initially), and showcase peers' successes to build confidence. Offer training and recognition: e.g., an "AI-Enhanced Teaching" certification for faculty who upskill. Also make clear that AI isn't replacing instructors or advisors – it's a tool (and prove it by tracking that staff roles remain critical, just improved). Over time, as comfort grows, resistance should lessen – especially if early adopters report better outcomes and less burnout, others will follow.
5. **Integration and Technical Complexity:** There's risk of technical difficulties in making all these systems talk (especially given current fragmentation). Delays or failures here could hamper the whole project (e.g., if data integration isn't done, the AI can't function effectively). *Mitigation:* Invest in strong technical project management and perhaps external expertise for our architecture. Use iterative integration – don't wait for perfect integration to start AI use; do targeted integrations for each use case and gradually unify. Also consider interim solutions (manual data uploads) while automating pipelines in parallel. Build a robust testing plan for each integration to avoid critical system outages. Essentially, de-risk technically by not doing "big bang" replacements – keep core systems running while layering AI on top gradually.

Additionally, monitor emerging legal/regulatory risks (e.g., if government issues strict rules on AI in education or data export, we need to comply swiftly – our governance committee should track this).

First Investments ("What to do on Monday morning" – next 6–12 months)

To kickstart this transformation, we should immediately:

- **1. Establish Leadership & Governance:** Formally set up the University AI Office and name a high-level lead (Chief AI Officer or similar) to drive this. Empower a cross-functional AI governance committee to start meeting within weeks. This sends a signal of commitment and provides a structure for decisions.
- **2. Invest in Data Integration:** Begin building the unified data platform. Likely hire a data engineering team or allocate vendor budget now. This includes cloud infrastructure setup, licenses (if using tools), and starting to pipe in SIS and LMS data. It's foundational, so we cannot delay – aim to have a basic data lake and warehouse environment within 3–4 months.
- **3. Quick-Win AI Solutions:** Deploy a chatbot for Admissions/Student FAQs as a pilot (relatively low effort with today's tech). The investment is in a platform subscription or consulting to implement Q&A pairs. This will immediately benefit prospective students (improved service) and free staff time. Similarly, get a license or open-source tool for automated grading (like Gradescope) for a large course next semester – relatively small cost, high faculty visibility of AI helping.
- **4. Faculty & Staff Training Fund:** Allocate resources for an AI training program. For example, partner with a known center (maybe bring in an expert from a leading AI university to give workshops, or use online courses) to ensure our people start building AI competency. Also create small grants or incentives for instructors who pilot AI in their teaching (this encourages grassroots innovation to complement top-down).
- **5. Pilot Predictive Analytics in Advising:** Assemble a small data analytics task force (could start with existing Institutional Research staff plus an external data scientist consultant) to develop a prototype student risk model. Use last 2–3 years of data to flag factors and test accuracy. Acquire a visualization tool if needed for advisors to see the results. This pilot will get people used to data-driven decision support and we can refine it for broader use.
- **6. Strengthen Infrastructure & Security:** Ensure our IT network and systems are ready for increased AI and data load. That might mean upgrading storage, improving single sign-on integration, and importantly, security audits (AI systems can introduce new vulnerabilities). It might also include negotiating data-sharing agreements early (with cloud providers or edtech vendors) to ensure compliance and clarity on data ownership [55] .
- **7. Communication & Change Narrative:** Right away, start a communication campaign with internal stakeholders. Explain the vision (maybe an email or town hall from the President: "We're embarking on this AI transformation, here's why it's exciting for our mission and how it will help each of you"). Also, set up channels for input – e.g., an AI Transformation webpage or forum where progress is shared and people can ask questions. This transparency will build trust and enthusiasm.
- **8. Budget for Talent and Partnerships:** Identify key hires (maybe an AI architect, a data privacy officer dedicated to this, etc.) and start recruitment – the talent market in AI is competitive, so earlier is better. Similarly, earmark budget for strategic partnerships (maybe sign an MoU with an AI research lab for advisory support, or with a vendor for a favorable pilot pricing).
- **9. Policy Review Launch:** Task a small group (legal, academic affairs, student rep, AI ethicist if any) to start reviewing and drafting needed policy updates (e.g., academic integrity, data consent forms, IP policies for AI-generated content). These typically require committee approvals so starting now will ensure by the time tech is in place, policies are too.
- **10. Set Measurable Targets:** Decide on a few key KPIs to track from day one – e.g., student satisfaction with support services, retention rates, faculty time spent on admin tasks (from surveys),

etc. Baseline them now, so we can quantify improvement. This also helps focus everyone on outcomes, not just tech for tech's sake.

By focusing investments in these areas now, we create the conditions for success. We'll have leadership, infrastructure, quick wins, and a supportive culture building concurrently. This multi-pronged start is necessary given the 36-month horizon – we can't afford a year of planning with no action. These moves will demonstrate momentum within the first 6–12 months, which is critical to maintain support (Board and community) and to learn and iterate early.

In conclusion, this strategy sets us on a path to become an AI-native university that leverages the latest technology to enhance learning and operations while upholding our educational values. Through a balanced, well-governed approach, we can achieve substantial gains in student success, institutional efficiency, and innovative capacity. The transformation will require ongoing commitment – to data, to ethics, to training our people – but the reward is a smarter, more responsive university poised to lead in the digital era.

This comprehensive plan is grounded in current evidence and best practices from global peers. All claims and solutions are backed by research and case studies (as referenced), and we have built in evaluation checkpoints to remain adaptive. By taking these steps, our university can confidently navigate the AI revolution in higher education and emerge stronger, more inclusive, and future-ready. [19] [64]

Critical Assumptions & Validation

Throughout this strategy, we've made assumptions that should be validated to ensure success:

1. **Assumption:** Students will positively engage with AI tools (tutors, chatbots) and not reject them.
 Validation: Student focus groups and pilots in year 1. Survey users of the chatbot for satisfaction. Measure usage stats – if only a few use it, assumption fails and we need to adjust (perhaps marketing or tool design).
2. **Assumption:** Sufficient quality data exists to drive accurate predictive models (for retention, etc.).
 Validation: Conduct data audit and a pilot model. If initial model accuracy is very low, that indicates data gaps or noise – we'd then invest in better data collection (e.g., more frequent academic updates or a student survey to capture non-cognitive factors).
3. **Assumption:** Faculty grading time can be significantly reduced without sacrificing quality using AI assistance.
 Validation: A/B test in a course – one section uses AI-assisted grading, one traditional. Compare grading time logged and student outcomes/feedback on fairness. If results show minimal time saved or student concerns about fairness, refine the AI approach or training.
4. **Assumption:** AI can improve retention by enabling earlier, more targeted interventions.
 Validation: Monitor a cohort with the AI early alert system vs previous cohorts. If retention doesn't improve (considering other factors constant), examine if interventions are actually happening or effective – maybe need to adjust threshold or the human follow-up process.
5. **Assumption:** The data platform and AI architecture can be implemented with available resources (budget, talent).
 Validation: Make a detailed project plan with vendors/IT, get estimates. If it turns out we severely under-budgeted or lack skill, we may need to scale back initial scope or secure additional funding/partners.

6. **Assumption:** Students will adhere to AI usage policies if clearly taught (i.e., not massively cheating).
 Validation: Track academic integrity cases related to AI each term after policy rollout. If cases spike, then assume failed – need more education or technical controls.
7. **Assumption:** Vendor solutions (Canvas AI features, etc.) will work for our context and integrate.
 Validation: Pilot with vendor support. For example, turn on Canvas AI for a few courses, gather feedback from those instructors and any technical issues. If integration problems or features not useful, consider alternative approaches.
8. **Assumption:** The benefits (e.g., time savings, improved metrics) will outweigh costs and effort.
 Validation: Maintain a running cost-benefit analysis. After each major deployment, quantify resource invested vs. gain (like reduction in inquiry calls, etc.). If an initiative consistently consumes effort without visible benefit, be ready to pivot or drop it.
9. **Assumption:** AI models can be kept unbiased and accurate for our diverse student body.
 Validation: Conduct bias and accuracy tests as planned [5]. For instance, test the advising chatbot with queries from different demographic personas. If biases found, adjust training data or logic and retest until acceptable.
10. **Assumption:** Leadership and stakeholders will remain supportive through the changes (even if some failures occur).
 Validation: Frequent status reports to leadership with transparent discussion of wins and challenges. Gauge their sentiments in meetings. If support waning (maybe due to a public concern or slower progress), need to reinforce communication of incremental successes and possibly adjust expectations or project scope to regain confidence.

By explicitly validating these assumptions, we turn potential unknowns into monitored parameters, ensuring we can course-correct promptly.

Potential Failure Modes and Mitigations

Despite best efforts, things can go wrong. Here are five failure modes and how we'd address them:

1. **Failure Mode:** *"Data swamp"* – We ingest lots of data but it's inconsistent or not effectively used, leading models to fail or produce nonsense.
 Mitigation: Start with a strong data governance practice: define data owners for each source, clean data before use, and build incremental models (don't throw every variable in blindly). If initial models are poor, do a feature importance analysis to identify which data might be misleading. Possibly simplify – it's better to have a few high-quality features than a messy many. Also, invest in an enterprise data dictionary so everyone understands the data.
2. **Failure Mode:** *Bias or ethical scandal* – e.g., an AI advising tool gives a student bad advice that goes viral, or a predictive model is accused of discrimination.
 Mitigation: Respond swiftly: pause the tool if needed, conduct an investigation (with outside experts if credibility needed), and be transparent about findings and fixes [68]. From the start, have an ethics review process that might catch such issues. Additionally, have a clear disclaimer on tools that they are decision support, not infallible. If a particular model (say admissions AI) is too sensitive, decide to not deploy until thoroughly vetted or stick to human process there. The key is to act responsibly and openly to maintain trust.
3. **Failure Mode:** *Low Adoption – "The AI tools are there but no one uses them."* Perhaps faculty turn off the AI features, or students ignore the chatbot.
 Mitigation: Diagnose reasons: usability issues? Lack of awareness? Trust? Then address specifically. For faculty, maybe the AI suggestions need to be more customizable or integrated into their workflow; gather their feedback and iterate. For students, maybe do a campaign: e.g., orientation includes how to use the AI assistant effectively, or

even embed some required interactions (like a freshman scavenger hunt where they must ask the chatbot something). If after serious effort a tool still isn't used, maybe it's solving the wrong problem – be ready to pivot resources to more valued areas.

4. **Failure Mode:** *Cost Overrun* – We might find costs (cloud compute for AI, vendor licenses, hiring data scientists) much higher than anticipated, risking budget issues.
 Mitigation: Implement cost monitoring from day one. Many AI services have usage-based pricing – set budgets and alerts. Optimize: e.g., schedule heavy jobs for overnight when cheaper, or use open-source models on our own servers for frequent tasks if that's cost-effective. If overrun is unavoidable, prepare a case for additional investment by showing the ROI (tie it to outcomes to justify the spend). Also, prioritize use cases: ensure the ones that save or bring value are scaled first so that those gains can offset other costs.
5. **Failure Mode:** *Technical collapse at critical time* – e.g., AI system goes down during registration or gives incorrect grades at scale, causing chaos.
 Mitigation: Plan redundancy and fail-safes. Always have a manual fallback for critical processes (if AI fails during registration, staff temporarily step in or revert to manual process). For grading, ensure instructors can override and sanity-check before release. Do not fully remove humans from loops until system is proven extremely reliable. Conduct stress testing in non-critical periods. Also, communicate to users that new tech might have hiccups and we have backup plans – so one glitch doesn't completely erode confidence. If a collapse happens, respond proactively: communicate transparently to affected users, fix quickly, and analyze root cause to prevent recurrence.

The overarching mitigation strategy is to pilot, monitor, and have layered defenses (technology, process, human oversight) such that a single failure doesn't cascade.

Pilot Experiments for the Next 4–8 Weeks

To build momentum and learn quickly, we propose six focused pilot experiments to run in the next 1–2 months:

Pilot 1: "Smart Advising FAQ Bot"
- **Hypothesis:** A fine-tuned AI chatbot (using GPT-4 or similar) can correctly answer at least 80% of common advising questions (course requirements, deadlines) and reduce human email queries by 50%.
- **Cohort:** New incoming students (e.g., 100 volunteer freshmen) during course registration period.
- **Method:** Deploy a chatbot via our portal or Facebook Messenger. Seed it with advising handbook content and Q&A. Students are told to ask it anything they'd normally email an advisor. Compare volume of advising emails from this cohort vs a control group of freshmen not using the bot.
- **KPI:** % of questions the bot answers correctly (verified via student feedback or advisor review logs); reduction in advisor email load; student satisfaction rating of answers.
- **Guardrails:** Advisors double-check a random sample of bot responses daily (to intercept any wrong info). The bot carries a disclaimer it's beta. Any question it's unsure about, it advises contacting an advisor (to avoid confident wrong answers).
- **Duration:** 4 weeks (covering initial registration + add/drop). At end, survey participants and advisors. If results are good, we'll iterate on content and consider scaling to all students.

Pilot 2: "AI-Assisted Grading in Intro Computing"
- **Hypothesis:** Using Gradescope's AI grouping or an LLM to draft feedback for programming assignments will cut TA grading time by at least 30% while maintaining grading quality.
- **Cohort:** One section of CS101 (~50 students) this semester. Another section serves as control (traditional

grading).
- **Method:** Train TAs on using the tool. For one assignment or two, TAs use AI to cluster similar wrong answers and apply batch grades, also have AI suggest feedback phrases. Control section TAs do usual process. Compare time logs (we'll ask TAs to record how long they spend) and consistency of grades (maybe have an instructor double-mark a sample from both groups to see if variation differs).
- **KPI:** Grading turnaround time, TA time spent, and student feedback on feedback quality (short survey: "Did you find the feedback useful?" without telling them which section had AI). Also measure if grade distributions are similar (ensuring fairness).
- **Guardrails:** Instructor reviews AI-suggested grades for a subset before release to avoid any major errors. Communicate to students that an AI may assist feedback but instructor oversees it, to prevent confusion.
- **Duration:** 6 weeks (through two assignments). Then evaluate and gather TA input: did it actually help or create overhead?

Pilot 3: "Retention Early Alert in Faculty Portal"
- **Hypothesis:** Providing faculty with an AI-generated "at-risk" indicator for their students (based on LMS activity and early grades) within the first 4 weeks will lead to more interventions (faculty reach-outs) and better midterm performance of those students.
- **Cohort:** Perhaps 10 volunteer instructors teaching first-year courses (~300 students total).
- **Method:** Develop a simple model (could even be rule-based to start: e.g., no login in 7 days + <60% on first quiz = flag) to identify at-risk students by week 4. Show these flags on the faculty's LMS dashboard or via email. Faculty are asked to do some outreach (email or talk to student) and log it. Meanwhile, similar courses without flags serve as comparison. Track whether flagged students subsequently submit assignments and their midterm grades vs non-flagged.
- **KPI:** Faculty intervention rate (did they actually reach out for X% of flags), subsequent assignment submission rate of flagged students vs historically, midterm score improvements relative to expectation. Faculty attitudes also measured via quick feedback ("were the flags useful?").
- **Guardrails:** Emphasize this is a supplemental tool, not an official label – to avoid bias, perhaps don't show risk score, just a nudge "consider checking in with Student A." Make sure flagged info isn't visible to students or others (to avoid stigma).
- **Duration:** 8 weeks (first half of term). Evaluate outcomes at midterm. If promising (e.g., majority of flagged who were contacted improved attendance or performance), we refine model and plan to expand.

Pilot 4: "AI Office Hours Assistant" (Supporting faculty during office hours or Q&A forums)
- **Hypothesis:** An AI assistant like "Jill Watson-lite" on a discussion forum can handle at least 30% of student questions with correct answers, reducing response time significantly (especially after hours).
- **Cohort:** One online course or a large class with heavy Q&A (maybe an Economics 101 with lots of questions). ~100 students.
- **Method:** Use an LLM connected to the course textbook and FAQ. Enable it on the forum with a bot account that attempts to answer routine questions ("When is homework due?" or "Can you explain concept X?") We'll focus on factual or clarifying questions, not subjective ones. The instructor/TA monitors.
- **KPI:** Fraction of student questions answered by the AI before a human intervenes; average response time for those (likely minutes) vs non-AI responses in control classes (maybe hours); accuracy of answers (instructor rates them). Also student feedback – did they find it helpful?
- **Guardrails:** Limit AI to answer only if confidence high (we can instruct it or use an approval mechanism). TAs must review daily to correct any mistakes publicly. The AI identifies itself clearly ("I'm a virtual TA").
- **Duration:** 4 weeks mid-semester. Evaluate logs and feedback. If accuracy is, say, >90% on fact-based questions and students appreciated quick replies, we consider scaling up.

Pilot 5: “Personalized Module in Math”
- **Hypothesis:** An adaptive learning module for a specific math topic (e.g., algebra basics) will yield higher mastery (post-test scores) for students who use it vs those who only get traditional instruction.
- **Cohort:** 50 students in a remedial math workshop, randomly split into two groups: one uses the AI module (like an ALEKS or Khan Academy adaptive set of exercises) for 2 weeks on that topic, the other studies via textbook and static problem set.
- **Method:** Pre-test both groups on the topic to gauge baseline. Group A goes through adaptive software (which adjusts difficulty, provides hints via AI tutor), Group B does normal practice with teacher support. Then post-test.
- **KPI:** Improvement from pre- to post-test between groups; also track time spent – maybe adaptive group achieves equal or better in less time (efficiency). Student feedback on experience (was it engaging, or frustrating?).
- **Guardrails:** Ensure both groups have access to help (Group B from teacher, Group A from the AI hints and some teacher oversight if needed). Ethical: because it’s learning intervention, make sure no group is left at disadvantage – since it’s short, if one method underperforms significantly, we can give those students extra tutoring afterward.
- **Duration:** ~3 weeks (including testing). This informs how effective such tools are for our students; if great, push more adaptive tech, if not, investigate why (content issue, tech issue, etc.).

Pilot 6: “AI-Enhanced Student Support Center”
- **Hypothesis:** AI can triage and respond to basic inquiries in the student support center (IT helpdesk or library help), reducing human ticket volume and speeding resolution for simple issues by at least 40%.
- **Cohort:** The IT helpdesk – for a period, have an AI chat pop-up or IVR that attempts to handle top 5 common issues (password reset instructions, Wi-Fi setup, etc.). Or do similarly for library FAQs (how to find resources).
- **Method:** Deploy a Q&A bot on the IT support page. Log how many issues it resolves (user doesn’t escalate) vs how many get escalated to human. Compare with historical data on those issues.
- **KPI:** Resolution rate by AI, average time to resolution (immediate vs waiting for human reply), reduction in tickets reaching staff. Staff feedback too – did it meaningfully lighten their load?
- **Guardrails:** Make it easy to “talk to a human” if AI fails. Ensure the AI responses are validated (we feed it with correct solutions for known issues). Start with only very predictable issues so risk of wrong advice is low.
- **Duration:** 4-8 weeks. If we see, for example, AI resolved 200 queries that otherwise would have been tickets, that’s a success and we can formalize it. If users still bypass it mostly, maybe they need a different interface or trust building.

Each pilot is scoped small (4-8 weeks) with clear hypotheses and measurements, so we can learn quickly and build confidence. They also each target different stakeholders (students, faculty, staff) to get broad insights. We will document results and use them to refine the larger roll-out plan. Early wins will be publicized internally to maintain momentum, and any early missteps will be quietly corrected before scaling. This experimental mindset is part of becoming an AI-native, learning organization [74] – we continuously test and improve.

---

[1] [2] Cognii - Blog - Florida International University Partners With Cognii to Implement AI in Online Education
https://www.cognii.com/blog/o26727-Florida-International-University-Partners-With-Cognii-to-Implement-AI-in-Online-Education

[3] [4] [71] Squirrel Ai Learning Sets Guinness World Record for "The most users to take an online mathematics lesson in 24 hours"
https://www.prnewswire.com/news-releases/squirrel-ai-learning-sets-guinness-world-record-for-the-most-users-to-take-an-online-mathematics-lesson-in-24-hours-302324623.html

[5] [6] [7] [35] [52] [53] Georgia Tech’s Jill Watson Outperforms ChatGPT in Real Classrooms | College of Computing
https://www.cc.gatech.edu/news/georgia-techs-jill-watson-outperforms-chatgpt-real-classrooms

[8] [9] [10] Instructors Will Now See AI Throughout a Widely Used Course Software
https://www.chronicle.com/article/instructors-will-now-see-ai-throughout-a-widely-used-course-software

[11] [12] [13] [14] [15] [16] [54] [70] A Year of AI in Learning Management Systems
https://www.govtech.com/education/higher-ed/a-year-of-ai-in-learning-management-systems

[17] [18] [20] [26] [33] [51] [63] [69] AI in Higher Education: Top Use Cases You Need To Know
https://smartdev.com/ai-use-cases-in-higher-education/

[19] [68] Ellucian's AI Survey of Higher Education Professionals Reveals Surge in AI Adoption Despite Concerns Around Privacy and Bias | Ellucian
https://www.ellucian.com/newsroom/ellucians-ai-survey-higher-education-professionals-reveals-surge-ai-adoption-despite

[21] [22] [23] [37] [50] Using AI Chatbots to Freeze ‘Summer Melt’ in Higher Ed -- Campus Technology
https://campustechnology.com/articles/2017/03/07/using-ai-chatbots-to-freeze-summer-melt-in-higher-ed.aspx

[24] [25] [43] [44] AI & Academic Integrity | Center for Teaching Innovation
https://teaching.cornell.edu/generative-artificial-intelligence/ai-academic-integrity

[27] [28] [29] [30] [31] [47] Deakin's Genie: a virtual digital assistant out of the bottle | Deakin
https://www.deakin.edu.au/about-deakin/news-and-media-releases/articles/deakins-genie-a-virtual-digital-assistant-out-of-the-bottle

[32] [41] [42] [45] [46] [55] [61] [62] [64] Governance of Generative AI in Higher Education: Lessons From the Top 50 Universities in the United States - Freedom For All Americans
https://freedomforallamericans.org/generative-ai-governance-higher-education/

[34] [36] [48] [49] [59] [60] [65] How the AI for Teaching & Learning Framework on AWS is transforming the student and teacher experience | AWS Public Sector Blog
https://aws.amazon.com/blogs/publicsector/how-the-ai-for-teaching-learning-framework-on-aws-is-transforming-the-student-and-teacher-experience/

[38] [39] [72] How Higher Ed Institutions Are Using Built-In Generative AI Tools | EdTech Magazine
https://edtechmagazine.com/higher/article/2025/06/how-higher-ed-institutions-are-using-built-generative-ai-tools

[40] [73] Most students outrunning faculty in AI use, study finds
https://www.insidehighered.com/news/tech-innovation/artificial-intelligence/2023/10/31/most-students-outrunning-faculty-ai-use

[56] Can a ‘Family of Bots' Reshape College Teaching? | EdSurge News
https://www.edsurge.com/news/2018-04-13-can-a-family-of-bots-reshape-college-teaching

[57] [58] Applying AI in higher education – An inevitable trend for the future | Nhan Dan Online
https://en.nhandan.vn/applying-ai-in-higher-education-an-inevitable-trend-for-the-future-post155207.html

[66] [67] On the Ethics of A/B Testing - CITP Blog
https://blog.citp.princeton.edu/2014/07/08/on-the-ethics-of-ab-testing/

74 Understanding Ethical Considerations in A/B Testing - Stellar
https://www.gostellar.app/blog/ethical-considerations-in-ab-testing